

# When Shared Memory Hurts: Stale-Coordination Failure in Multi-Agent Particle Collection under a Rotating Latent Field

Anonymous Author(s)

Submission Id: «OpenReview submission id»

## Abstract

Sharing observations across a team of agents improves coordination under a stationary transport field — but what happens when the field drifts? A team that pools a long observation history can converge confidently on a stale consensus, driving all collectors toward the wrong direction while independent agents, sampling noisier but more diverse estimates, collectively cover more ground. This paper asks where exactly that reversal occurs and what governs it.

Working in a stochastic particle-collection task with a confirmed stationary coordination benefit ( $\bar{\Delta} = +1.19$  particles,  $n=32$  matched seeds), we introduce a linearly rotating field and measure the last memory window  $L_{\max}$  that maintains a statistically significant benefit, across an order-of-magnitude range of rotation speeds and eight window lengths.

The central finding is evidence for a *stale-coordination failure mechanism*: the shared arm reduces angular noise but leaves angular bias unchanged, so beyond a critical window length it drives all collectors confidently toward a systematically wrong direction, reducing the inter-agent angular diversity that makes coordination beneficial. At both interior speeds ( $\omega \in \{1.5, 5.0\}$ ), instrumented runs show the mechanism’s predicted directional signatures: U-shaped shared-arm angular error profiles with observed minima at their respective  $L_{\max}$  (paired  $t$ -tests:  $t(31) \geq 3.40$ ,  $p < 0.001$  at both speeds), and inter-agent angular dispersion structurally zero in the shared arm and non-zero in the independent arm. The coordination benefit landscape under rotation is *regime-dependent*: qualitatively different outcomes arise at different rotation speeds. Four regimes emerge from this study. A *linear-rotation regime* ( $L_{\max}/T_{\text{corr}} \approx 0.30$ , sinc attenuation negligible) exhibits a clean significance-based crossover at  $L_{\max} \approx 0.30 T_{\text{corr}}$  at two interior speeds; a *sinc-attenuation regime* ( $\omega \in \{2.0, 2.5, 3.0\}$ ) dampens the estimate amplitude across the tested  $L$  range, suppressing the sharp crossover; a *rapid-rotation regime* ( $\omega=10.0$ ) drives the predicted  $L_{\max}$  below integer resolution; and a *long-episode regime* ( $\omega=1.0$ ,  $T_{\text{corr}} \approx H$ ) exhibits an unexplained mechanism — raising the question of what governs which regime a given  $(\omega, H)$  pair falls into.

## CCS Concepts

• **Computing methodologies** → **Multi-agent planning**; *Cooperative and collaborative learning*.

## Keywords

multi-agent coordination; non-stationary environments; memory window; communication benefit; stochastic particle collection; rotating latent field

## ACM Reference Format:

Anonymous Author(s). 2027. When Shared Memory Hurts: Stale-Coordination Failure in Multi-Agent Particle Collection under a Rotating Latent Field. In *Proc. of the 26th International Conference on Autonomous Agents and Multiagent Systems (AAMAS 2027), Hanoi, Vietnam, 3–7 May 2027*, IFAAMAS, 11 pages.

## 1 Introduction

Communication protocols for multi-agent systems are typically designed and evaluated against a fixed environment. In deployment, environments drift. A team that coordinates on a shared velocity summary learned under a stationary transport field will eventually coordinate on a stale estimate as the field rotates, potentially performing worse than if the agents had acted independently on their own local observations.

This paper asks a precise question: how long must an observation memory window be to maintain a positive coordination benefit when the latent transport field rotates at angular rate  $\omega$ ? We study this in the stochastic particle-collection task of SPS-C03 [1], where four mobile collectors must exploit a weak latent drift  $\alpha=0.06$  to collect more of  $N=256$  diffusing particles than a capacity-matched independent team with no shared information. The confirmed stationary baseline is  $\bar{\Delta}=+1.19$  particles (SD 2.44,  $p < 0.001$ , 32 matched seeds).

We extend this baseline along the temporal axis by introducing a linearly rotating field  $\theta(t) = \theta(0) + \omega \cdot t \cdot dt$  and asking two implementation-level questions: (i) what memory window  $L$  should agents use, and (ii) how does the best achievable  $L$  scale with the rotation rate  $\omega$ ?

**Main result.** At two interior rotation speeds ( $T_{\text{corr}} \in \{33, 10\}$ ), the pre-registered significance-based threshold satisfies

$$L_{\max}(\omega)/T_{\text{corr}} \approx 0.30, \quad T_{\text{corr}} = \frac{1}{\omega \cdot dt},$$

where  $T_{\text{corr}}$  is the one-radian decorrelation timescale, and  $L_{\max}$  is the significance-based threshold (Section 3.2; two related thresholds,  $L^*$  and the coordination-collapse boundary, are defined there). At these two speeds the empirical ratio  $L_{\max}/T_{\text{corr}} = 0.30$  holds exactly; the window spans 0.27 rad (slow) and 0.20 rad (mid) of total field rotation at the threshold under the corrected  $(L-1)$  lag formula — the *linear-rotation regime* where sinc attenuation is negligible and the midpoint-lag approximation holds. Other tested speeds fall into qualitatively different regimes (see Section 5.2). The constant  $c \approx 0.30$  rests on two interior anchors; the slow anchor ( $p=0.049$ ) was confirmed by a pre-specified sensitivity analysis at  $n=64$  ( $t=+1.87$ ,



This work is licensed under a Creative Commons Attribution International 4.0 License.

$p < 0.05$ ). In the slow condition, a window beyond  $L_{\max}$  reverses the benefit:  $\bar{\Delta}$  falls from +0.94 at  $L=10$  to  $-0.38$  at  $L=20$  ( $\omega=1.5$  rad/step). At mid ( $\omega=5.0$ ), the mean  $\bar{\Delta}$  remains positive until  $L^* \in (30, 45)$ , far beyond  $L_{\max}=3$ ; a long window can degrade but does not necessarily reverse the benefit within the tested range.

**Why memory eventually hurts.** The mechanism is not a degradation of each agent’s individual directional accuracy. Both arms accumulate the same angular lag ( $\omega \cdot (L-1) \cdot dt/2$  radians at window midpoint); the shared arm merely reduces angular *noise* while leaving the bias unchanged. Once the window spans enough rotation, the shared arm drives all  $M$  collectors confidently toward the same stale direction, eliminating the coverage diversity that makes independent estimates collectively useful. This is a policy nonlinearity: the capacity-matching controller maps confident-but-wrong directions into globally poor coverage.

### Contributions.

- (1) We provide direct measurement evidence for a stale-coordination failure mechanism: at both interior speeds, instrumented runs show U-shaped shared-arm angular error with observed minima at their respective  $L_{\max}$ , and inter-agent directional dispersion that is structurally zero in the shared arm while the independent arm retains diversity (Figure 3). These directional signatures are consistent with the proposed mechanism. We do not measure spatial trajectory overlap, spatial coverage directly, or provide a formal mediation analysis connecting angular dispersion to captures; those remain open.
- (2) We identify a four-regime taxonomy of the coordination benefit landscape under rotation: (i) a linear-rotation regime ( $\omega L_{\max} dt \approx 0.30$  rad) with a clean crossover at  $L_{\max}/T_{\text{corr}} \approx 0.30$ ; (ii) a sinc-attenuation regime where amplitude damping suppresses the sharp crossover; (iii) a rapid-rotation regime where  $L_{\max}$  falls below integer resolution; and (iv) a long-episode regime ( $T_{\text{corr}} \approx H$ ) with an unresolved mechanism. The  $c \approx 0.30$  result is the quantitative finding within regime (i).
- (3) We show that the empirical crossover occurs at lower  $L_{\max}$  than the bias/noise-parity prediction ( $L^* \propto (\omega dt)^{-2/3}$ ), a different functional form, attributable to policy nonlinearity.
- (4) We show that exponential decay does not follow the same  $L_{\max} \propto T_{\text{corr}}$  pattern at the two interior speeds, making the sliding window a more tractable choice in the studied regime.
- (5) We provide a pre-registered, reproducible experiment with event-keyed matched counterfactuals and an explicit reproduction gate anchored to SPS-C03.

## 2 Related Work

*Non-stationary multi-agent environments.* Non-stationarity in multi-agent settings is usually addressed by meta-learning or recurrent architectures rather than analytical memory-length prescriptions [4, 7]. These approaches adapt  $L$  implicitly through learned gates but offer no closed-form guidance for practitioners. Patiño et al. [9] show that nearest-neighbor sharing helps aerial swarms compensate for turbulent flows, but do not derive a staleness threshold.

*Windowed and discounted estimation.* Sliding-window UCB (SW-UCB) [5] and discounted UCB [8] apply the same bias/variance logic to non-stationary bandit problems: a window of length  $L \approx \sqrt{n/\Delta}$  is near-optimal under a known breakpoint rate  $\Delta$ . Our sliding-window controller is the cooperative analogue of SW-UCB applied to field-direction estimation across agents. We find a pattern consistent with  $1/\omega$  scaling at two interior rotation speeds ( $L_{\max}/T_{\text{corr}} \approx 0.30$ ), steeper than the bias/noise-parity prediction ( $L^* \propto \omega^{-2/3}$ ); this pattern is preliminary and requires further validation.

*Communication under information constraints.* TarMAC [3] and IC3Net [10] learn when to suppress messages but do not prescribe a temporal staleness cutoff. The information-bottleneck approach of Wang et al. [12] optimises message bandwidth but not memory depth. Our contribution is orthogonal: we hold the message fixed (a count-weighted team mean velocity) and vary only the observation history depth.

*Particle collection.* Wang et al. [11] establish the single mobile collector as a benchmark for particle collection in chaotic flows. Our multi-agent extension isolates whether shared summaries lower the minimum signal strength needed to exploit a latent field, and the current paper asks how robustly that benefit persists under field rotation. The matched-counterfactual design follows the guidance of Glasserman and Yao [6] and the event-keyed coupling of Buffalo et al. [2].

## 3 Model and Theory

### 3.1 Particle-collection task

The arena is the unit square  $[0, 1]^2$  with reflecting boundaries.  $N=256$  passive particles diffuse with step size  $\sigma=0.06$  per  $dt=0.02$  second and are advected by a latent field with strength  $\alpha=0.06$ . Four collectors (sensing radius  $r=0.16$ , max speed  $v_{\max}=0.12$ ) run for  $H=67$  steps; a particle is permanently captured when it enters a collector’s disc. The estimand is

$$\Delta_s = Y_{\text{shared},s} - Y_{\text{indep},s},$$

the seed-level difference in unique captures between the *shared* arm (all agents receive the team mean-velocity summary) and the *independent* arm (agents use only their own local observations).

*Matched counterfactual coupling.* Both arms run from the *identical* pre-generated noise tensor  $W_s \in \mathbb{R}^{H \times N \times 2}$ , initial particle positions  $X_s^0$ , initial collector positions  $C_s^0$ , and deterministic field sequence  $\theta_s(t)$ ; only the message channel differs. This event-keyed coupling [2] ensures that  $\Delta_s$  is a within-seed contrast: the stochastic environment is held constant while communication alone varies. Equation (1) is evaluated over the shared team history (shared arm) or each agent’s own history (independent arm); all other aspects of the simulation are identical.

### 3.2 Rotating-field model

The latent field direction evolves deterministically:

$$\theta(t) = \theta(0) + \omega \cdot t \cdot dt, \quad \theta(0) \sim \text{Uniform}[0, 2\pi).$$

At  $\omega=0$  the task reduces to SPS-C03 exactly. The field rotates deterministically, so its cosine similarity with the current direction

oscillates rather than decaying; nonetheless the one-radian decorrelation timescale—the lag at which cosine similarity first drops to  $\cos(1) \approx 0.54$ —is a natural staleness timescale:

$$T_{\text{corr}} = \frac{1}{\omega \cdot dt} \quad (\text{steps}).$$

Note the  $dt$  factor: it is the *per-step* rotation  $\omega \cdot dt$ , not the rate  $\omega$  itself, that determines how quickly observations become stale.

**Threshold definitions.** Three distinct  $L$ -based thresholds appear in the results; they coincide only under specific conditions:

- **Significance threshold** ( $L_{\text{max}}$ ): the largest  $L$  at which the pre-registered one-sided  $t$ -test gives  $p < 0.05$ . This is a statistical-detectability threshold, not a behavioural reversal.
- **Zero-crossing threshold** ( $L^*$ ): the value of  $L$  at which  $\mathbb{E}[\Delta]$  changes sign, estimated by interpolation between adjacent grid points.
- **Coordination-collapse boundary**: the  $L$  at which the expected captures under the shared arm fall below the independent arm — equivalent to the zero-crossing. We use this phrase only when  $L_{\text{max}} \approx L^*$ , i.e. when the significance threshold and the sign reversal are in the same  $L$ -interval.

### 3.3 Memory controllers

**Sliding-window controller (window).** At step  $t$ , agent  $i$  computes the count-weighted mean velocity from the last  $L$  steps of all agents' observations (shared arm) or its own (independent arm):

$$\hat{v}(t, L) = \frac{\sum_{\tau=\max(0, t-L+1)}^t \sum_i n_i(\tau) \bar{v}_i(\tau)}{\sum_{\tau} \sum_i n_i(\tau)}, \quad (1)$$

where  $n_i(\tau)$  is the count of particles observed by agent  $i$  at step  $\tau$  and  $\bar{v}_i(\tau)$  is their local mean velocity. At  $L=H$  this reduces to the full-episode SPS-C03 controller.

**Capacity-matching action rule.** The estimated velocity  $\hat{v}$  is converted to a movement command by the capacity-matching blend:

$$a_m = \text{unit}(\mathbf{w}_t \hat{\mathbf{f}} + (1 - \mathbf{w}_t) \hat{\mathbf{d}}_m), \quad \hat{\mathbf{f}} = -\text{unit}(\hat{\mathbf{v}}), \quad (2)$$

where  $\hat{\mathbf{d}}_m = \text{unit}(\bar{\mathbf{p}}_m)$  is the unit vector toward agent  $m$ 's local particle centroid,  $\bar{\mathbf{p}}_m$  is the mean relative position of visible particles,  $N_{\text{max}}$  is the per-agent sensing capacity (particles per step at full-disc occupancy), and the blend weight  $\mathbf{w}_t = \min(0.7, 2 \bar{n}_t / N_{\text{max}})$  rises with the current mean per-agent particle count  $\bar{n}_t$  (instantaneous, at step  $t$ ). When  $\hat{v}=0$  (no observations in the window), agents fall back to the density term alone. All agents in the shared arm receive the same  $\hat{v}$  and hence the same  $\hat{\mathbf{f}}$ ; they differ only in  $\hat{\mathbf{d}}_m$ . The policy nonlinearity arises here: a high blend weight  $\mathbf{w}_t$  causes all collectors to commit strongly to  $\hat{\mathbf{f}}$ , so a stale  $\hat{v}$  drives the entire team toward a single wrong direction, eliminating the diversity in  $\hat{\mathbf{d}}_m$  that would otherwise provide spatial coverage.

**Exponential-decay controller (decay).** The same sum with exponentially decreasing weights  $\lambda^{t-\tau}$ ,  $\lambda = \exp(-1/L)$ . The effective memory length equals  $L$  in expectation, but older observations are downweighted smoothly rather than discarded at a hard cutoff.

**Table 1: Angular bias and shared-arm noise at and beyond  $L_{\text{max}}$ . At  $L=L_{\text{max}}$  the bias is well below  $\sigma_{\theta, \text{shared}}$ ; the ratio first exceeds 1 at approximately  $2L_{\text{max}}$ , consistent with a policy-driven crossover at lower  $L$  than the directional-MSE model predicts.**

| $\omega$ ( $T_{\text{corr}}$ ) | $L$                      | $\omega(L-1)dt$ | $\beta$   | $\beta/\sigma_{\theta, \text{shared}}$ |
|--------------------------------|--------------------------|-----------------|-----------|----------------------------------------|
| 1.5 (33)                       | 10 ( $=L_{\text{max}}$ ) | 0.27 rad        | 0.135 rad | 0.34                                   |
| 1.5 (33)                       | 20                       | 0.57 rad        | 0.285 rad | 1.02                                   |
| 5.0 (10)                       | 3 ( $=L_{\text{max}}$ )  | 0.20 rad        | 0.100 rad | 0.14                                   |
| 5.0 (10)                       | 10                       | 0.90 rad        | 0.450 rad | 1.14                                   |

( $\alpha=0.06$ ,  $\sigma=0.06$ ,  $M=4$ ,  $dt=0.02$ ;  $\bar{K} \approx 8$  is the episode-mean per-agent particle count defined in Section 3.4)

### 3.4 Signal-to-noise analysis and its limits

A natural first hypothesis is that sharing  $M$  agents' observations improves the effective sample size by  $M$ , giving a signal-to-noise ratio  $\text{SNR}_{\text{shared}}/\text{SNR}_{\text{indep}} = \sqrt{M}$  independent of  $L$  and  $\omega$ . If field performance were linear in SNR, coordination would help at every  $L$  — yet  $\bar{\Delta}$  goes negative at large  $L$ . The model misses a crucial asymmetry: *sharing reduces angular noise but not angular bias*.

Define the angular bias (lag to the window midpoint),

$$\beta(L, \omega) = \frac{\omega \cdot (L-1) \cdot dt}{2} \quad (\text{radians}),$$

and the angular noise of the shared arm,

$$\sigma_{\theta, \text{shared}}(L) = \frac{\sigma}{\alpha \sqrt{M \bar{K} L dt}},$$

where  $\bar{K} = \mathbb{E}_t[\bar{n}_t]$  is the mean per-agent particle count averaged over steps (the episode mean of the instantaneous count  $\bar{n}_t$  used in the action rule;  $\bar{K} \approx 8$  at the studied parameters). Both arms carry the same bias  $\beta$ ; sharing only reduces  $\sigma_{\theta}$  by  $\sqrt{M}$ . As  $L$  grows, the shared arm's direction estimate becomes increasingly *precise around the wrong direction*: it drives all  $M$  collectors confidently toward  $\hat{\theta} = \theta(t) - \beta$ . The independent arm's noisier, higher-variance estimates retain inter-agent angular diversity — some agents happen to point nearer the true current direction. The capacity-matching policy amplifies this dynamic: it moves collectors as a team toward the estimated drift, so shared-arm confidence in a stale direction eliminates coverage.

The bias/noise-parity threshold ( $\beta \approx \sigma_{\theta, \text{shared}}$ ) gives  $L^* \propto (\omega dt)^{-2/3}$ . This is a different functional form from the empirical  $L_{\text{max}} \propto (\omega dt)^{-1}$ : the ratio  $L^*/T_{\text{corr}}$  scales as  $(\omega dt)^{1/3}$  and is therefore not a universal constant. Because the actual threshold is set by the *policy's* nonlinear sensitivity rather than directional MSE, the empirical crossover is earlier and follows a steeper scaling.

**Regime structure of the window estimate.** The midpoint-lag formula  $\hat{\theta} = \theta(t) - \omega(L-1)dt/2$  (and the bias/noise analysis above) is valid only when  $\omega L dt \ll 1$ , i.e. when the window spans a small total rotation. Beyond this linear-rotation regime, the window estimate is governed by the exact geometric-sum formula in Appendix C, whose amplitude scales as  $|\sin(\omega L dt/2)|/(L |\sin(\omega dt/2)|)$ . Two qualitatively distinct failure modes result:

- (1) *Sinc-attenuation regime*. When  $\omega L dt$  is substantial but the amplitude is still positive, both the noise-reduction benefit and the stale-direction cost are scaled down together. The mechanism operates but weakly, producing gradual degradation rather than a sharp crossover.
- (2) *Rapid-rotation regime*. When  $T_{\text{corr}}$  is so short that  $c T_{\text{corr}} < 1$  step (predicted  $L_{\text{max}} < 1$ ), no integer-valued  $L$  lies inside the linear-rotation regime. Even  $L=1$  carries a per-step rotation  $\omega dt$  as its lag angle, and the resulting benefit is too small to detect at practical sample sizes.

A third anomaly at very slow rotation ( $\omega=1.0$ ,  $T_{\text{corr}} \approx H$ ) lies outside both accounts; its mechanism is not resolved in this work. Section 5.2 characterises all four regimes empirically.

## 4 Experimental Setup

### 4.1 Grid design

We test four non-stationary rotation speeds and one stationary baseline, plus a supplementary third-anchor condition ( $\omega=2.5$ , between slow and mid), eight  $L$  levels, and two methods:

| Label      | $\omega$ (rad/step) | $\omega \cdot dt$ | $T_{\text{corr}}$ |
|------------|---------------------|-------------------|-------------------|
| stationary | 0                   | 0                 | $\infty$          |
| very_slow  | 0.75                | 0.015             | 67                |
| slow       | 1.5                 | 0.030             | 33                |
| supp.      | 2.5                 | 0.050             | 20                |
| mid        | 5.0                 | 0.100             | 10                |
| fast       | 17.0                | 0.340             | 3                 |

$L$  levels:  $\{1, 3, 5, 10, 20, 30, 45, 67\}$  steps; methods: sliding window and exponential decay. The primary design yields 40 active conditions (4 non-stationary  $\omega \times 5$  primary  $L \times 2$  methods), plus 24 intermediate- $L$  cells for slow and mid, plus stationary- $\omega$  cells, and 16 supplementary cells for the  $\omega=2.5$  anchor (8  $L$  levels  $\times$  1 method – window only). We call slow and mid the *interior* rotation speeds: they are the two conditions where  $L_{\text{max}}$  falls strictly within the grid range  $(1, H)$  and the linear-rotation regime ( $\omega L_{\text{max}} dt \lesssim 0.30$  rad) applies. Intermediate  $L$  values were concentrated at these interior speeds (slow, mid,  $\omega=2.5$ ) by design: very\_slow has  $L_{\text{max}}=67$  (boundary, unresolvable within the grid) and fast has predicted  $L_{\text{max}} < 1$  step (below grid resolution). Post-hoc exploratory runs at  $\omega=1.0$  ( $T_{\text{corr}}=50$ ) and  $\omega=2.0$  ( $T_{\text{corr}}=25$ ), each 8  $L$  levels  $\times$  32 seeds (window method), are reported in Sec. 5.2.

### 4.2 Matched counterfactuals

Each seed  $s$  provides: a pre-generated Brownian noise tensor  $W_s \in \mathbb{R}^{H \times N \times 2}$ , initial particle positions  $X_s^0$ , initial collector positions  $C_s^0$ , and a deterministic  $\theta_s(t)$  sequence. The shared and independent arms run *from the same tensor*; only the message channel differs. This follows the event-keyed coupling of Buffalo et al. [2].

### 4.3 Seeds and statistical power

Seeds 9001–9032 (32 per cell) provide  $\text{SE} \approx 0.45$  per cell ( $\text{SD} \approx 2.5$ ), sufficient to detect  $|\bar{\Delta}| \approx 0.9$  at 80% power one-sided. The stationary condition uses seeds 6001–6128 (128 per cell) to resolve the smaller  $L$ -dependent effect at  $\omega=0$ .

### 4.4 Reproduction gate

Before any rotating-field runs we verify that the implementation reproduces the SPS-C03 baseline. Gate criterion:  $\omega=0$ ,  $L=1$  (stateless window), seeds 6001–6032:  $\bar{\Delta} \in [+0.69, +1.69]$ , sign count  $\geq 20/32$ .

**Gate result:**  $\bar{\Delta}=+1.188$  ( $\text{SE} = 0.213$ ), sign=20/32. One-sided  $t$ -test (same test used for all  $L_{\text{max}}$  cells):  $t=5.58$ ,  $df=31$ ,  $p < 0.001$ . Both gate criteria satisfied. ✓

### 4.5 Pre-registered analysis

Primary: one-sided  $t$ -test per cell ( $H_0 : \Delta \leq 0$ ,  $\alpha_{\text{test}}=0.05$ ).  $L_{\text{max}}(\omega)$  is the largest  $L$  with  $p < 0.05$  for the window method. Scaling fit: force-through-origin OLS of  $L_{\text{max}}$  on  $T_{\text{corr}}$  using all rotation speeds with defined  $L_{\text{max}}$ .

## 5 Results

Recall the three threshold types defined in Section 3.2: significance threshold ( $L_{\text{max}}$ ), zero-crossing ( $L^*$ ), and coordination-collapse boundary. Applied to the current data: for the slow condition ( $\omega=1.5$ ),  $L_{\text{max}}=10$  and  $L^* \in (10, 20)$  are in the same interval, making  $L_{\text{max}}$  an approximate coordination-collapse boundary. For the mid condition ( $\omega=5.0$ ),  $L_{\text{max}}=3$  but  $L^* \in (30, 45)$ :  $L_{\text{max}}$  marks only statistical detectability; mid does not exhibit a coordination-collapse boundary within the tested grid.

### 5.1 Per-cell coordination benefit

Table 2 shows  $\bar{\Delta}$ , SE, and significance for all window-method cells. The stationary condition ( $\omega=0$ ) gives significant benefit at  $L \in \{1, 3, 10\}$  but not at  $L=30$ , consistent with the stationary controller’s own optimal window (pooling over a full episode loses little given  $\bar{\Delta}$  is nearly flat at  $L \leq 10$ ).

For  $\omega=\text{slow}$  ( $T_{\text{corr}}=33$ ) the coordination benefit is large and positive for  $L \leq 10$  and sharply negative for  $L \geq 20$ :  $\bar{\Delta}$  goes from  $+0.94$  ( $L=10$ ,  $p=0.049$ ) to  $-0.38$  ( $L=20$ ,  $p=0.79$ ), a swing of  $> 1.3$  particles when the window doubles from 10 to 20 steps.

For  $\omega=\text{mid}$  ( $T_{\text{corr}}=10$ ) significance is lost earlier but the character differs from slow:  $\bar{\Delta}=+0.91$  at  $L=3$  ( $p=0.004$ ) falls to  $+0.22$  at  $L=5$  (NS) and the estimated mean remains positive through  $L=30$ . The mean does not turn negative until  $L=45$  ( $\bar{\Delta}=-0.25$ ). The mid condition therefore primarily demonstrates *loss of detectable advantage* at  $L > 3$ , not the sharp sign reversal seen in slow.

For  $\omega=\text{fast}$  ( $T_{\text{corr}}=3$ ) no  $L$  value achieves consistent significance. The isolated  $p=0.023$  at  $L=5$  is an outlier: adjacent cells  $L \in \{1, 3, 10\}$  are all non-significant. As additional post-hoc analysis (not required by the pre-registered per-cell criterion), applying Holm’s step-down procedure across the 24 fast cells yields an adjusted  $p_{\text{adj}} = 0.023 \times 24 = 0.55$  for the minimum raw  $p$ , confirming no fast cell survives even relaxed correction. Importantly, the theoretical prediction for this condition is  $L_{\text{max}} \approx 0.30 \times 3 = 0.9$  steps, which falls *below* the minimum grid value of  $L=1$ . The null result is therefore expected: the  $L$  grid cannot resolve  $L_{\text{max}} < 1$  step, and the fast condition provides no independent evidence against the scaling law.

A supplementary condition  $\omega=2.5$  ( $T_{\text{corr}}=20$  steps, between slow and mid) was run with 32 seeds across all 8  $L$  levels to probe the scaling at an intermediate rotation speed.  $\bar{\Delta}$  remains positive at every  $L$  (range  $+0.69$  to  $+1.72$ ), with a local minimum of  $+0.78$  at  $L=10$  ( $t=1.91$ ). The one-sided  $t$ -test is significant ( $p < 0.05$ ) at every

$L$  through  $L=45$  ( $t=1.86$ ) but not at  $L=67$  ( $t=1.67$ ); no crossover to negative  $\bar{\Delta}$  was observed. The interpretation of  $L_{\max}$  for this condition is discussed in Section 5.2.

## 5.2 $L_{\max}$ versus $T_{\text{corr}}$

Table 3 summarises  $L_{\max}$  (last significant  $L$  under the pre-registered per-cell  $p < 0.05$  criterion, window method) per rotation speed and the ratio  $L_{\max}/T_{\text{corr}}$ . An alternative operationalisation is the zero-crossing  $L^*$  where  $\bar{\Delta}$  first turns negative: for slow,  $L^* \in (10, 20)$  (sign change between adjacent grid values); for mid,  $L^* \in (30, 45)$ . At slow,  $L_{\max}$  and  $L^*$  approximately agree (both locate the crossover near  $L=10-20$ ): the significance threshold and the zero-crossing are close, consistent with a genuine coordination-collapse boundary in that range. At mid, however,  $L_{\max}=3$  while  $L^* \in (30, 45)$  — a tenfold gap.  $\bar{\Delta}$  remains positive at  $+0.22$  through  $L=30$  and only turns negative at  $L=45$ ;  $L_{\max}=3$  therefore marks *loss of statistical detectability* rather than a coordination-collapse boundary. Under the zero-crossing operationalisation, mid gives  $c^* \in (3.0, 4.5)$  — an order of magnitude above the significance-based  $c=0.30$ . The  $c=0.30$  result is therefore primarily anchored on the slow condition where the two operationalisations agree; mid corroborates the  $1/T_{\text{corr}}$  scaling pattern for  $L_{\max}$  but measures a different crossover quantity at the same  $L_{\max}$  value. The significance-based  $L_{\max}$  is used throughout as the pre-registered quantity.

The two interior rotation speeds (slow and mid) give exactly  $L_{\max}/T_{\text{corr}} = 0.30$ . The very slow condition ( $T_{\text{corr}}=67$ ) gives  $L_{\max}=67$  — the full episode — which is consistent with the  $1/\omega$  scaling but does not sharpen the constant because any true  $L_{\max}$  between 45 and 67 would produce the same reported value within our  $L$  grid. The fast condition ( $T_{\text{corr}}=3$ ) has no reliable  $L_{\max}$ .

The supplementary condition  $\omega=2.5$  ( $T_{\text{corr}}=20$ , predicted  $L_{\max} \approx 6$ ) was run with 32 seeds to test whether  $c=0.30$  holds at an intermediate rotation speed. Applying the same pre-registered criterion used for slow and mid (last  $L$  with one-sided  $t$ -test  $p < 0.05$ ,  $n=32$ ), we obtain  $L_{\max}(\omega=2.5)=45$  ( $t=1.86$ ,  $p < 0.05$ );  $L=67$  falls below threshold ( $t=1.67$ ,  $p > 0.05$ ). This gives  $c=45/20=2.25 - 7.5\times$  above the two-anchor value. However, the  $t$ -profile is *non-monotone*: after peaking at  $L=1$  ( $t=3.07$ ) it dips at  $L=10$  ( $t=1.91$ ), recovers at  $L=20$  ( $t=2.38$ ) and  $L=30$  ( $t=1.97$ ), then declines through  $L=45$  ( $t=1.86$ ) to  $L=67$  ( $t=1.67$ ). A genuine sharp crossover — as at slow, where  $\bar{\Delta}$  drops from  $+0.94$  to  $-0.38$  across a single  $L$  doubling — produces a monotone decline and a clear sign reversal; the  $\omega=2.5$  profile shows neither. The  $\omega=2.5$  condition therefore *cannot serve as a clean third anchor*: the stopping criterion delivers a technical  $L_{\max}=45$ , but not a genuine sharp crossover. **The headline  $c \approx 0.30$  continues to rest on the two interior anchor points (slow, mid).** We report it as a preliminary value; the  $\omega=2.5$  result suggests that the true relationship may be more complex — in particular, that the crossover sharpness itself depends on  $\omega$ .

An exploratory run at  $\omega=2.0$  ( $T_{\text{corr}}=25$ , predicted  $L_{\max} \approx 7.5$ ) with 32 seeds (window method) reveals a second broad-benefit anomaly: the one-sided  $t$ -test is significant at every  $L$  in the grid, with  $t$ -statistics ranging from  $t=5.57$  at  $L=1$  to a minimum of  $t=1.93$  at  $L=45$ , and recovering to  $t=2.07$  at  $L=67$  ( $p < 0.05$  throughout). The full-episode  $L=67$  window remains beneficial, giving  $L_{\max} \geq 67$

( $c \geq 67/25=2.68$ ) — a lower bound. Like  $\omega=2.5$ , this condition does not yield a genuine crossover and cannot serve as an anchor.

An exploratory run at  $\omega=10.0$  ( $T_{\text{corr}}=5$ , predicted  $L_{\max} \approx 1.5$  under  $c=0.30$ ) with 32 seeds reveals the *rapid-rotation regime*. The only nominally significant cell is  $L=30$  ( $t=2.04$ ,  $p < 0.05$ ), but this result should not be taken as a real coordination benefit. At  $L=30$  the window estimate amplitude — computed from the geometric-sum formula in Appendix C — is only 4.7% of the single-step amplitude ( $\approx 0.047$  from the geometric-sum formula): near-zero cancellation, the window spanning nearly one full rotation ( $\omega L dt=6.0 \text{ rad} \approx 2\pi$ ). The estimate direction is  $-166^\circ$  relative to the current field — nearly backwards — with negligible amplitude. An estimate of this quality cannot explain a genuine coordination benefit; the  $t=2.04$  result is consistent with the expected 0.4 false positives from testing eight  $L$  cells at  $\alpha=0.05$ . The  $L$  values flanking the anomaly confirm the diagnosis: at  $L=1$  (the predicted  $L_{\max}$  at  $c=0.30$ ), the estimate is accurate (magnitude 1.0,  $0^\circ$  lag), yet the benefit is not significant ( $t=1.27$ ,  $p > 0.10$ ) — consistent with a genuine small effect that is underpowered at  $n=32$ . At  $L=45$  and  $L=67$  the benefit is negative ( $\bar{\Delta} < 0$ ), confirming the mechanism still operates; the failure is that  $L_{\max}$  is below integer resolution, not that the mechanism is absent. Table 4 shows the full  $t$ -profile; the spike at  $L=30$  is visually isolated and the two adjacent cells ( $L=45$ ,  $L=67$ ) are both negative.

**Sinc-attenuation and long-episode regimes.** A unified account of the anomalous conditions emerges from the sinc-attenuation factor derived in Appendix C: at each  $L$ , the shared arm’s direction estimate is amplitude-scaled by  $\text{sinc}(\omega L dt/2)$ . At the two clean interior anchors,  $L_{\max}$  coincides with  $\omega L_{\max} dt/2 = 0.15 \text{ rad}$ , giving  $\text{sinc} \approx 0.996$  — negligible attenuation, so the “confident-wrong-direction” mechanism operates at essentially full signal strength at the crossover.

For  $\omega \in \{2.0, 2.5, 3.0\}$  the picture is different. At  $L=45$  the sinc factor is 0.87 ( $\omega=2.0$ ), 0.80 ( $\omega=2.5$ ), and 0.72 ( $\omega=3.0$ ); at  $L=67$  it falls to 0.73, 0.59, and 0.45 respectively. This amplitude attenuation weakens the mechanism throughout the upper  $L$  range: both the beneficial noise-reduction effect and the harmful stale-direction effect are scaled down together, so the net coordination benefit degrades gradually rather than collapsing sharply. Three independent runs confirm this sinc-attenuation regime: at  $\omega=3.0$  ( $n=64$ ,  $T_{\text{corr}}=16.7$ , predicted  $L_{\max} \approx 5$ ), the  $t$ -profile is non-monotone —  $L=5$  is significant ( $t=2.57$ ) as the linear-regime formula predicts, but benefit re-emerges and persists through  $L=67$  ( $t \in [1.77, 3.32]$ ), with  $\bar{\Delta}$  remaining positive ( $> 0.6$ ) at all large  $L$ . This is precisely the sinc-attenuation signature: the estimate is wrong-directed at large  $L$  but amplitude-damped enough that coverage diversity is partially preserved.

The  $\omega=1.0$  anomaly does not fit the sinc story:  $\text{sinc} \approx 0.999$  at  $L=3$  (where benefit collapses), and 0.927 at  $L=H$  — both mild. The elevated per-seed SD at  $\omega=1.0$  ( $\approx 3.8$  particles vs. 3.2 at slow) and the positive-but-underpowered benefit at  $L \geq 3$  ( $\bar{\Delta}=+0.28$ ,  $n=32$ ) are consistent with high seed-to-seed variability driven by the episode spanning  $\omega H dt=1.34 \text{ rad}$ : enough rotation that the episode-level outcome depends strongly on the initial field angle  $\theta(0)$ . Whether  $L_{\max}=1$  is a genuine mechanistic threshold or a power artefact is not resolved at  $n=32$ .

**Table 2: Window-method per-cell results.**  $n=32$  seeds per cell (non-stationary  $\omega$ );  $n=128$  for stationary  $L \in \{1, 3, 10, 30\}$ .  $\bar{\Delta}$  = mean shared-independent captures; (\*) =  $p < 0.05$  one-sided  $t$ -test. Dashes: intermediate  $L$  not run for that  $\omega$ . †Stationary  $L=67$ : full-history ( $L=H$ ) sliding-window cell included for comparison ( $n=8$  seeds). The SPS-C03 baseline ( $\bar{\Delta}=+1.19$ ) is reproduced by the  $\omega=0$ ,  $L=1$  reproduction gate (Section 4, App. B); the  $L=67$  stationary cell coincidentally yields the same value; field stationarity removes temporal bias, reducing but not eliminating the effect of window length on estimator variance.  $\omega=2.5$ : supplementary condition; see Sec. 5.2. Exploratory conditions  $\omega \in \{1.0, 2.0, 3.0, 10.0\}$  are not included here; their  $L_{\max}$  values appear in Table 3 and their  $t$ -profiles are discussed in Sec. 5.2.

| $\omega$ ( $T_{\text{corr}}$ ) | $L=1$  | $L=3$  | $L=5$  | $L=10$ | $L=20$ | $L=30$ | $L=45$ | $L=67$  |
|--------------------------------|--------|--------|--------|--------|--------|--------|--------|---------|
| 0 ( $\infty$ )                 | +1.19* | +0.81* | —      | +0.72* | —      | +0.25  | —      | +1.19*† |
| 0.75 (67)                      | +0.53  | +0.47  | +0.06  | +0.22  | +0.50  | +0.41  | +0.66  | +0.78*  |
| 1.5 (33)                       | +1.47* | +1.50* | +1.66* | +0.94* | −0.38  | −0.28  | −0.47  | −0.44   |
| 2.5 (20)                       | +1.72* | +1.47* | +1.19* | +0.78* | +1.06* | +0.88* | +0.88* | +0.69   |
| 5.0 (10)                       | +0.88* | +0.91* | +0.22  | +0.25  | +0.56  | +0.16  | −0.25  | +0.03   |
| 17.0 (3)                       | +0.47  | +0.47  | +0.78  | −0.09  | +0.09  | −0.38  | +0.13  | +0.25   |

These observations suggest that  $\omega H dt$  — the total field rotation per episode — is a meaningful regime variable for the clean-crossover condition. Both confirmed clean anchors (slow:  $\omega H dt=2.01$  rad; mid: 6.70 rad) have their  $L_{\max}$  in the sinc-negligible regime ( $\omega L_{\max} dt/2 = 0.15$  rad). The three-speed sinc-attenuation confirmation ( $\omega \in \{2.0, 2.5, 3.0\}$ ) tightens the regime boundary: the linear-to-sinc-attenuation transition lies in  $\omega \in (3.0, 5.0)$ , i.e.  $T_{\text{corr}} \in (10, 16.7)$  steps. Varying  $H$  at fixed  $\omega$  would directly test whether a sharp crossover re-emerges when the grid can access  $\omega L_{\max} dt \approx 0.30$  rad without encountering substantial sinc attenuation.

**Crossover condition.** Both clean interior anchors satisfy

$$L_{\max}/T_{\text{corr}} = 0.30 \quad (\text{equivalently, } \omega \cdot L_{\max} \cdot dt = 0.30 \text{ rad}).$$

**Table 3:  $L_{\max}$  (last  $L$  with  $p < 0.05$ , window) and ratio to  $T_{\text{corr}}$ , with regime label. Bold: two clean interior points (slow, mid) in the linear-rotation regime with  $L_{\max}/T_{\text{corr}} = 0.30$ . ¶  $\omega=1.0$ : only  $L=1$  significant ( $t=2.68$ ); all  $L \geq 3$  non-significant ( $t < 0.5$ ). §  $\omega=2.5$ :  $L_{\max}=45$  by criterion ( $t=1.86$ ), but  $t$ -profile non-monotone (sinc-attenuation regime). ‡  $\omega=2.0$ : all  $L$  significant throughout;  $L_{\max}$  is a lower bound ( $\geq 67$ ; sinc-attenuation). †  $\omega=10.0$ : only  $L=30$  significant ( $t=2.04$ ), but estimate amplitude is 4.7% at  $L=30$  (near-zero cancellation); likely false positive — see text. ||  $\omega=3.0$ :  $n=64$ ;  $L=5$  significant ( $t=2.57$ ) as predicted, but benefit persists through  $L=67$  ( $t=2.86$ ); sinc-attenuation regime,  $L_{\max}$  is a lower bound.**

| $\omega$ | $T_{\text{corr}}$ | $L_{\max}$           | $\omega L_{\max} dt$ | $L_{\max}/T_{\text{corr}}$ | Regime        |
|----------|-------------------|----------------------|----------------------|----------------------------|---------------|
| 0.75     | 67                | 67                   | 1.00 rad             | 1.00                       | linear (bdy.) |
| 1.0      | 50                | 1¶                   | 0.02 rad             | 0.02                       | long-ep.      |
| 1.5      | 33                | 10                   | <b>0.30 rad</b>      | <b>0.30</b>                | <b>linear</b> |
| 2.0      | 25                | $\geq 67^{\ddagger}$ | $\geq 2.68$ rad      | $\geq 2.68$                | sinc-att.     |
| 2.5      | 20                | 45§                  | 2.25 rad             | 2.25                       | sinc-att.     |
| 3.0      | 17                | $\geq 67^{  }$       | $\geq 4.02$ rad      | $\geq 4.02$                | sinc-att.     |
| 5.0      | 10                | 3                    | <b>0.30 rad</b>      | <b>0.30</b>                | <b>linear</b> |
| 10.0     | 5                 | —†                   | —                    | —                          | rapid-rot.    |
| 17.0     | 3                 | —                    | —                    | —                          | rapid-rot.    |

**Table 4:  $\omega=10.0$  window-method per-cell results ( $n=32$ ). Amplitude of the  $L=30$  estimate: 4.7% (near-zero cancellation, direction  $-166^\circ$ ). The sole significant cell is consistent with the 0.4 false positives expected from eight  $\alpha=0.05$  tests;  $L=1$  (predicted  $L_{\max}$  under  $c=0.30$ ) is non-significant despite a high-quality estimate (amplitude 1.0,  $0^\circ$  lag).**


| $L$ | $\bar{\Delta}$ | SE   | $t$     |
|-----|----------------|------|---------|
| 1   | +0.56          | 0.44 | +1.27   |
| 3   | +0.22          | 0.39 | +0.56   |
| 5   | +0.50          | 0.47 | +1.06   |
| 10  | +0.41          | 0.45 | +0.91   |
| 20  | +0.44          | 0.42 | +1.03   |
| 30  | +1.03          | 0.51 | +2.04 * |
| 45  | −0.19          | 0.63 | −0.30   |
| 67  | −0.16          | 0.55 | −0.29   |

Under the corrected angular-lag formula  $\beta = \omega(L-1)dt/2$ , the direction lag at  $L_{\max}$  is 0.135 rad (slow) and 0.100 rad (mid) — approximately but not identically consistent. The total angular span of the window ( $\omega(L_{\max}-1)dt$ ) is 0.27 rad (slow) and 0.20 rad (mid): the previously reported “ $17^\circ$ ” was an artefact of using  $\omega L_{\max} dt$  rather than  $\omega(L_{\max}-1)dt$  for the window span. The  $L_{\max}/T_{\text{corr}}=0.30$  ratio remains the correct empirical summary. The  $\omega=2.0$  and  $\omega=2.5$  conditions do not yield reliable  $L_{\max}$  estimates and are excluded from this calculation.

Figure 1 plots  $L_{\max}$  versus  $T_{\text{corr}}$ , including the two clean interior anchors, the boundary very\_slow point, and the three exploratory conditions ( $\omega \in \{1.0, 2.0, 2.5\}$ , all excluded from the fit; see caption). The force-through-origin fit on the two interior anchors plus the boundary-censored very\_slow point gives  $L_{\max} \approx 0.85 T_{\text{corr}}$  ( $R^2=0.80$ ), inflated by censoring; the two interior anchors alone give  $c=0.30$  exactly.

### 5.3 Crossover is sharp at $\omega = \text{slow}$

Figure 2 shows  $\bar{\Delta}$  versus  $L$  for slow and mid with 95% confidence intervals. At slow,  $L_{\max}=10$  and  $L^* \in (10, 20)$  co-locate the crossover



**Figure 1:**  $L_{\max}$  (last significant  $L$ , window method) vs.  $T_{\text{corr}} = 1/(\omega \cdot dt)$ . Filled circles: two clean interior anchors (slow, mid) on the  $c=0.30$  line; open square: very\_slow (boundary, consistent with  $c=0.30$  but censored); open upward triangles:  $\omega=2.5$  and  $\omega=2.0$  (anomalously high  $L_{\max}$ , excluded from fit); open downward triangle:  $\omega=1.0$  (anomalously low  $L_{\max}=1$ , excluded from fit). Dashed:  $L=T_{\text{corr}}$  reference; solid: force-through-origin fit on two interior anchors plus boundary point ( $c=0.85$ , inflated by censoring); dotted:  $c=0.30$  through the two interior anchors.


in the same  $L$ -interval, confirming slow as a genuine coordination-collapse boundary (Section 3.2). At mid the sharp-crossover character is absent:  $\bar{\Delta}$  shrinks gradually rather than collapsing, and post-hoc power for the  $L=5$  cell ( $\bar{\Delta}=+0.22$ ,  $d \approx 0.09$ ) is below 15% — consistent with the detectability interpretation of  $L_{\max}=3$  established in Section 5.2.

#### 5.4 Directional diagnostics of the proposed mechanism

The coordination-performance crossover is consistent with the stale-coordination failure mechanism, but does not isolate it: a lower  $L_{\max}$  could also arise from particle depletion or agent crowding. To test the mechanism directly, we ran instrumented episodes at both slow ( $\omega=1.5$ ) and mid ( $\omega=5.0$ ), each with  $n=32$  seeds, recording at every step  $t$ : the true field direction  $\theta(t) = \theta_0 + \omega t dt$ ; the shared arm’s direction estimate  $\hat{\theta}_{\text{shared}}$  (windowed team mean, identical for all agents by construction); and each agent’s independent estimate  $\hat{\theta}_{\text{indep}}^{(m)}$  (own windowed mean). At mid we used  $L \in \{1, 3, 10, 30, 45\}$ ; at slow, all eight  $L$  levels.

Figure 3 shows both conditions. The same two signatures appear at each speed:

- (1) **Shared-arm angular error is U-shaped with observed minimum at  $L_{\max}$  at both speeds.** Slow ( $\omega=1.5$ ): mean angular error falls from 0.768 rad at  $L=1$  to 0.286 rad at  $L=10$  ( $L_{\max}$ ), then rises to 0.575 rad at  $L=67$ . Adjacent paired comparisons (same 32 seeds across  $L$  levels,  $r \approx 0.80$  within-seed) confirm the minimum:  $E(5) > E(10)$ : paired  $t(31)=8.40$ ,



**Figure 2:**  $\bar{\Delta}$  (mean shared – independent captures,  $\pm 95\%$  CI) vs. memory window  $L$  for slow ( $T_{\text{corr}}=33$ , blue) and mid ( $T_{\text{corr}}=10$ , orange), window method,  $n=32$  seeds. Filled markers:  $p < 0.05$ ; open markers: not significant. Dashed verticals mark  $L_{\max}$  for each speed. The crossover for slow is sharp:  $+0.94$  at  $L=10$  to  $-0.38$  at  $L=20$ , a swing of  $>1.3$  particles.

$p < 0.001$ , bootstrap 95% CI  $[+0.053, +0.085]$  rad;  $E(20) > E(10)$ : paired  $t(31)=3.40$ ,  $p=0.0009$ , 95% CI  $[+0.016, +0.058]$  rad. Mid ( $\omega=5.0$ ): mean shared-arm angular error decreases from 0.770 rad at  $L=1$  to 0.471 rad at  $L=3$  ( $L_{\max}$ ), with a paired difference of 0.299 rad ( $t(31)=34.40$ ,  $p < 0.001$ ), before increasing to 0.538 rad at  $L=10$  and 1.215 rad at  $L=30$  as sinc attenuation dominates beyond the linear-rotation regime. The observed profile is U-shaped, with its observed minimum at  $L_{\max}$ . Additional paired confirmation of the right arm ( $r \approx 0.51$ ):  $E(10) > E(3)$ : paired  $t(31)=5.30$ ,  $p < 0.001$ , 95% CI  $[+0.041, +0.090]$  rad. The independent arm has no corresponding minimum at either speed.

- (2) **Shared arm inter-agent dispersion is structurally zero.**


Because all agents receive the identical team estimate, directional spread is zero for every  $L$  at both speeds. Independent agents maintain 0.585 rad of diversity at slow  $L=10$  and 0.928 rad at mid  $L=3$ , with gradual reduction as  $L$  grows.

The directional ingredients predicted by the mechanism are directly measured at both interior speeds; their causal mediation of capture performance is not isolated. What changes at  $L_{\max}$  is not a collapse of shared-arm dispersion — that is structural at every  $L$  — but the interaction: zero diversity compounded by a stale field direction rather than an accurate one.

#### 5.5 Pooled coordination benefit across $L$ levels

Table 5 reports  $\bar{\Delta}$  pooled across all  $L$  levels per  $\omega$  (window method, seed-level aggregation: for each seed,  $\bar{\Delta}_s$  is the mean  $\Delta$  across all  $L$  levels; the  $t$ -test runs on 32 seed means). The prior version of this table treated seed- $L$  pairs as independent (inflating  $n$  to 256–512); the corrected  $n=32$  uses clustered SEs. Only the stationary condition achieves  $p < 0.05$ ; the non-stationary conditions are directionally





**Figure 3: Directional diagnostics at slow ( $\omega=1.5$ , top) and mid ( $\omega=5.0$ , bottom);  $n=32$  seeds (9001–9032) each. Shaded bands:  $\pm 1$  SE. Left panels: mean angular error  $|\hat{\theta} - \theta_{\text{true}}|$ . Shared arm (blue) shows U-shaped angular error with observed minimum at  $L_{\text{max}}$  at both speeds. Slow ( $L_{\text{max}}=10$ ): confirmed by adjacent paired  $t$ -tests ( $t(31) \in \{8.40, 3.40\}$ , both  $p < 0.001$ ;  $B=10,000$  bootstrap CIs exclude zero). Mid ( $L_{\text{max}}=3$ ): error decreases from 0.770 rad at  $L=1$  to 0.471 rad at  $L_{\text{max}}=3$  (paired  $t(31)=34.40$ ,  $p < 0.001$ ), then rises to 0.538 rad at  $L=10$  and 1.215 rad at  $L=30$ ; right-arm confirmation:  $E(10) > E(3)$ , paired  $t(31)=5.30$ ,  $p < 0.001$ . Beyond  $L_{\text{max}}$  the shared arm rises due to angular lag (slow) or sinc attenuation (mid,  $\omega L dt$  large). Right panels: inter-agent directional dispersion. Shared arm is zero by construction for all  $L$ ; independent arm retains diversity (0.585 rad at slow  $L=10$ ; 0.928 rad at mid  $L=3$ ). Zero shared-arm dispersion is structural at every  $L$ ; the harmful regime is zero diversity combined with a stale direction, not a collapse of diversity at  $L_{\text{max}}$ .**

consistent but individually underpowered at the pooled level. The Q3 hypothesis (monotone decrease with  $\omega$ ) is not supported: pooled  $\bar{\Delta}$  is non-monotone (0.74, 0.45, 0.50, 0.34, 0.22 for stationary through fast). This is expected because pooling over all  $L$  averages short- $L$  cells (where benefit is robust) with long- $L$  cells (where it degrades); the degradation is  $L$ -conditional, not pooled.

## 5.6 Sliding window versus exponential decay

Table 6 shows  $L_{\text{max}}$  for both methods. The decay controller does not follow the  $L_{\text{max}} \propto T_{\text{corr}}$  scaling: it achieves  $L_{\text{max}}=67$  at very\_slow (NS,  $L_{\text{max}}$  undefined),  $L_{\text{max}}=5$  at slow, and  $L_{\text{max}}=10$  at mid — an irregular pattern with no consistent constant  $c$ . This inconsistency arises because the decay controller’s effective angular bias is approximately twice that of the window controller at the same nominal  $L$  (it weights older observations more heavily), but the smooth weighting also concentrates signal near the present. The two effects partially offset each other in a  $L$ -dependent and  $\omega$ -dependent

**Table 5: Pooled  $\bar{\Delta}$  across all  $L$  levels, window method.  $n=32$  seeds per  $\omega$ ;  $\bar{\Delta}$  is the mean over per-seed  $L$ -averages;  $p$ -values are one-sided. Earlier version used  $n=\text{seed-}L$  pairs (256–512), inflating precision; corrected here.**

| $\omega$ ( $T_{\text{corr}}$ ) | $n$ | $\bar{\Delta}$ | SE    | $p$   |
|--------------------------------|-----|----------------|-------|-------|
| 0 ( $\infty$ )                 | 32  | +0.742         | 0.325 | 0.015 |
| 0.75 (67)                      | 32  | +0.453         | 0.352 | 0.104 |
| 1.5 (33)                       | 32  | +0.500         | 0.353 | 0.084 |
| 5.0 (10)                       | 32  | +0.344         | 0.277 | 0.112 |
| 17.0 (3)                       | 32  | +0.215         | 0.296 | 0.237 |

**Table 6:  $L_{\text{max}}$  for both memory methods. The window method follows  $L_{\text{max}} \approx 0.30 \cdot T_{\text{corr}}$ ; the decay method does not. “—” = no  $L$  achieves  $p < 0.05$ .**

| $\omega$ ( $T_{\text{corr}}$ ) | $L_{\text{max}}$ (window) | $L_{\text{max}}/T_{\text{corr}}$ | $L_{\text{max}}$ (decay) | $L_{\text{max}}/T_{\text{corr}}$ |
|--------------------------------|---------------------------|----------------------------------|--------------------------|----------------------------------|
| 0.75 (67)                      | 67                        | 1.00                             | —                        | —                                |
| 1.5 (33)                       | 10                        | 0.30                             | 5                        | 0.15                             |
| 5.0 (10)                       | 3                         | 0.30                             | 10                       | 1.00                             |
| 17.0 (3)                       | —                         | —                                | —                        | —                                |

way that resists a simple scaling prediction. Importantly, the stale-coordination failure mechanism is not specific to the window controller: for slow, the decay controller’s  $\bar{\Delta}$  also turns negative beyond  $L_{\text{max}}$  ( $\bar{\Delta} \in [-0.19, -0.34]$  at  $L \in \{20, 30, 45\}$ ,  $n=32$ ), confirming that shared confidence in a stale direction eliminates coverage diversity under exponential weighting as well. The irregular  $L_{\text{max}}/T_{\text{corr}}$  ratio reflects an unpredictable *threshold*, not an absent mechanism.

**Practical observation:** At the two interior speeds studied, the sliding-window controller gives a  $L_{\text{max}} \propto T_{\text{corr}}$  relationship; the exponential-decay controller does not show a consistent pattern across the same conditions. Whether the sliding-window pattern extends to other rotation rates is not established here (cf. the  $\omega=2.5$  anomaly, Section 5.2).

## 5.7 Theoretical prediction vs. empirical result

The pre-experiment bias/noise-parity analysis (Section 3.4) predicts  $L^* \propto (\omega dt)^{-2/3}$ , a different functional form from the observed  $L_{\text{max}} \propto (\omega dt)^{-1}$ . A comparison of  $c$ -values is therefore ill-defined across rotation speeds (the implied ratio  $L^*/T_{\text{corr}}$  varies as  $(\omega dt)^{1/3}$ , not a constant). Pointwise: the bias/noise threshold predicts crossover at  $\beta/\sigma_{\theta, \text{shared}} \approx 1$ , but Table 1 shows this ratio at the observed  $L_{\text{max}}$  is only 0.34 (slow) and 0.14 (mid) — the empirical significance threshold is earlier than the theory predicts. We label the cause “policy nonlinearity” (capacity-matching sends all collectors toward a single estimated direction, so the coordination benefit collapses well before the shared arm’s individual directional accuracy drops to parity), but this is a qualitative label, not a derived functional form. Why the policy sensitivity produces  $L_{\text{max}} \propto \omega^{-1}$  rather than the  $L^* \propto \omega^{-2/3}$  from bias/noise parity is an open problem: the exponent difference is observed but not formally modelled in this work.



## 6 Discussion

*Practical memory-window heuristic.* For the specific parameter regime studied here ( $\alpha=0.06$ ,  $\sigma=0.06$ ,  $M=4$ ,  $dt=0.02$ , sliding-window controller), setting  $L \lesssim 0.30/(\omega \cdot dt)$  avoids the stale-coordination failure demonstrated in this work. The constant  $c \approx 0.30$  is a two-anchor preliminary value; re-validation is needed before applying it to different agent counts, field strengths, or coordination protocols. If  $\omega$  is unknown, estimating it from the rate of change of the team velocity estimate and adjusting  $L$  online is a natural but unvalidated extension.

*Connection to SW-UCB.* SW-UCB sets  $L \propto 1/\Delta$  where  $\Delta$  is a breakpoint rate [5]. Our results are consistent with a cooperative analogue  $L \propto 1/(\omega \cdot dt)$  with  $c \approx 0.30$  in the studied setting. The constant is smaller than the single-agent SW-UCB optimum because multi-agent sharing amplifies the consequence of directional bias; whether this analogy holds under other coordination protocols or reward structures is an open question.

*The exponential-decay controller.* Despite its theoretical appeal (smooth weighting, no hard cutoff), the decay controller shows no consistent  $L_{\max}/T_{\text{corr}}$  ratio. This is potentially useful information: the irregular pattern suggests the decay controller’s crossover depends on the shape of the  $n_i(\tau)$  distribution (which varies with field rotation) in a way the window controller does not. The irregular ratio does not undermine the mechanism’s generality: the benefit reversal occurs in both controllers (Sec. 5.6); what the decay controller lacks is a predictable failure threshold. A detailed analysis of this interaction is left for future work.

*Anchor robustness.* The slow anchor ( $\bar{\Delta}=+0.94$ ,  $p=0.049$ ) is borderline: a single additional seed going the opposite direction would push  $p$  above 0.05 and shift the reported  $L_{\max}$  from 10 to 5 steps. Furthermore, applying Holm’s step-down procedure uniformly across slow’s eight cells—as was done post-hoc for the fast condition (Section 5)—the adjusted threshold at  $L=10$ ’s rank (4th-smallest  $p$  of 8) is  $0.05/5 = 0.01$ ; the borderline  $p=0.049$  fails this threshold, shifting  $L_{\max}(\text{slow})$  from 10 to 5 steps (ratio  $5/33 \approx 0.15$ ). The mid anchor ( $L=3$ ,  $p=0.004$ ) survives Holm correction (threshold  $0.05/8 = 0.00625$ ). Under strict uniform Holm the two anchors give  $c = 0.15$  and  $c = 0.30$  respectively — no longer a consistent constant. The pre-registered criterion ( $p < 0.05$  per cell, uncorrected) is the basis for the  $c \approx 0.30$  claim. At  $n=32$  and the observed effect ( $\bar{\Delta}=+0.94$ , SD 3.20,  $d \approx 0.29$ ), post-hoc power is  $\approx 50\%$  — below the 80% design target (which assumed  $|\bar{\Delta}| \geq 0.90$ , SD 2.5) — and the one-sided 95% lower confidence bound ( $\approx -0.03$  particles) is consistent with zero.


*Pre-specified sensitivity analysis ( $n=64$ ).* Motivated by the borderline  $p$ -value, we pre-specified a replication of 32 additional seeds (9033–9064) for the slow condition and pooled with the original 32. At  $L=10$  the pooled estimate is  $\bar{\Delta}=+0.70$  ( $t=+1.87$ ,  $df = 63$ ,  $p < 0.05$  one-sided,  $SE=0.38$ ), confirming the anchor after regression to the mean. The crossover to non-significance at  $L=20$  is preserved (pooled  $\bar{\Delta}=+0.14$ ,  $t=+0.37$ , NS). Effect shrinkage from  $+0.94$  to  $+0.70$  is expected for a borderline result; the positive direction and clean crossover at  $n=64$  provide the relevant robustness evidence.

*The regime taxonomy: post-hoc structure with mechanistic accounts.* The four-regime taxonomy was constructed from the data,

not pre-specified: the non-conforming speeds were observed first, then classified. A sceptical reader may therefore read the taxonomy as an explanation built around counterexamples rather than an independently tested prediction. We accept that framing; the taxonomy’s value is not as a prior but as a compact description of the observed diversity of outcomes, each with a mechanistic account. The non-conforming conditions ( $\omega \in \{1.0, 2.0, 2.5, 10.0\}$ ) are not random scatter — they have distinct, characterisable structure. The sinc-attenuation regime ( $\omega \in \{2.0, 2.5, 3.0\}$ , confirmed at three speeds) has a mechanistic account: amplitude damping weakens the “confident-wrong-direction” mechanism proportionally across the tested  $L$  range, producing broad benefit without a crossover. The regime boundary lies in  $\omega \in (3.0, 5.0)$  ( $T_{\text{corr}} \in (10, 16.7)$  steps), the tightest bound available from the current data. The rapid-rotation regime ( $\omega=10.0$ ) has a resolution-limit account:  $L_{\max}$  falls below the integer grid, the only significant cell ( $L=30$ ) is a false positive at 4.7% estimate amplitude, and the mechanism is verified to be present but undetectable at  $n=32$ . The long-episode regime ( $\omega=1.0$ ) remains an open problem. Together, these accounts give practitioners a testable diagnostic: compute  $\omega L dt$  at the candidate  $L$ . If this is  $\lesssim 0.30$  rad (linear-rotation regime),  $c \approx 0.30$  applies; if substantially larger, the regime is sinc-attenuated and the sharp crossover may not exist; if  $c T_{\text{corr}} < 1$ , the regime is rapid-rotation and integer-grid resolution is the limiting constraint. The central open question is: *what determines whether a given  $(\omega, H)$  pair produces a sharp coordination crossover?*

*Generality and functional dependence of  $c$ .* The  $c \approx 0.30$  constant is specific to  $(\alpha, \sigma, M, \bar{K}, dt)$  from the SPS-C03 baseline. A back-of-envelope argument for the dominant dependences: the angular noise of the shared arm scales as  $\sigma_{\theta, \text{shared}} \propto \sigma/(\alpha \sqrt{M \bar{K} L} dt)$ , so larger  $\alpha/\sigma$  or larger  $M$  reduces noise and allows the shared arm to remain beneficial to a longer lag before coverage loss dominates. This implies  $c$  is an increasing function of  $\alpha/\sigma$  and (weakly) of  $M$  — larger or stronger teams tolerate more staleness. We hypothesise that the scaling  $L_{\max} \propto T_{\text{corr}}$  extends to other memory-based cooperative controllers with a fixed-direction prior, but this has not been tested here; quantifying  $c(\alpha/\sigma, M)$  requires a parameter sweep outside the scope of this paper.

To probe whether the dimensionless scale  $\Gamma = \omega dt \cdot L$  is preserved at the *reward-optimal* window across  $\omega$ , we ran an exploratory  $\omega \times L$  sweep (7  $\omega$  values  $\times$  10  $L$  levels  $\times$  16 seeds, shared arm only) and defined  $L_{\text{opt}}(\omega) = \arg \max_L \mathbb{E}[Y_{\text{shared}}(L)]$  — the window maximising the shared arm’s own reward (not the coordination advantage  $\bar{\Delta}$ , and not the pre-registered threshold  $L_{\max}$ ). Figure 4 shows the result, with bootstrap 95% confidence intervals for  $\Gamma(\omega)$  obtained by resampling seeds ( $B=4000$ ). In the three linear-regime conditions ( $\omega \in \{0.5, 1.5, 5.0\}$ ,  $T_{\text{corr}} \in \{100, 33, 10\}$  steps), the point estimates are  $\omega \cdot dt \cdot L_{\text{opt}} \in \{0.30, 0.21, 0.30\}$  (Panel C); all three are consistent under a force-through-origin fit ( $\hat{\Gamma}=0.29$ , Panel B). However, the bootstrap CIs are wide — spanning roughly  $[0.05, 0.45]$ ,  $[0.03, 0.90]$ , and  $[0.20, 1.50]$  respectively — reflecting genuine argmax instability when selecting over ten discrete  $L$  values with only 16 seeds. The  $\omega=1.5$  point ( $\Gamma=0.21$ ) lies below the  $\pm 20\%$  band around 0.29, further illustrating that the three point estimates do not tightly cluster. The sweep therefore provides only weak, suggestive evidence that the



**Figure 4: Reward-optimal window sweep** ( $n=16$  seeds, shared arm only, exploratory).  $L_{\text{opt}}(\omega)$  denotes the  $L$  maximising shared-arm reward — distinct from  $L_{\text{max}}$  (last significant  $\Delta$ ) and from  $L^*$  (zero-crossing of  $\mathbb{E}[\Delta]$ ). **Left:** Shared-arm reward vs.  $L$  for  $\omega \in \{0.5, 0.75, 1.0, 1.5, 2.5, 5.0, 8.0\}$ . **Linear-regime curves** (solid blue) have reward peaks at moderate  $L$ ; **anomalous conditions** (dashed) are flat or boundary-limited. **Centre:**  $L_{\text{opt}}(\omega)$  vs.  $T_{\text{corr}} = 1/(\omega \cdot dt)$ . **Filled circles:** three linear-regime points; **dashed line:** force-through-origin fit, slope 0.29. **Open markers:** anomalous-regime points, excluded from fit. **Right:**  $\Gamma(\omega) = \omega \cdot dt \cdot L_{\text{opt}}$  vs.  $\omega$  with bootstrap 95% CIs ( $B=4000$  seed resamples). The  $\pm 20\%$  band is shown for reference; the  $\omega=1.5$  point lies below it, and all three linear-regime CIs are wide, reflecting argmax instability with  $n=16$  seeds over ten discrete  $L$  values. The result is suggestive but not confirmatory.

reward-optimal window obeys the same  $1/(\omega \cdot dt)$  scaling as the pre-registered threshold  $L_{\text{max}}$ , and the coincidence with  $c \approx 0.30$  should not be treated as corroborative. Non-linear-regime conditions deviate upward as expected: the long-episode anomaly ( $\omega=1.0$ ) gives  $\omega \cdot dt \cdot L_{\text{opt}}=0.9$ ; the sinc-attenuation condition ( $\omega=2.5$ ) gives 1.0; and the rapid-rotation condition ( $\omega=8.0$ ) gives a noise-dominated value.

**Limitations.** (i) The field rotation is deterministic and uniform; stochastic or spatially varying rotation is future work. (ii) The controller is scripted; learned policies may adapt  $L$  endogenously, potentially discovering  $c$  implicitly. (iii) We test one coordination mechanism (count-weighted team mean); other communication protocols may have different  $L_{\text{max}}$  profiles. (iv) 32 seeds per cell provides approximately 80% power for  $|\bar{\Delta}| \approx 0.9$ ; smaller effects at the boundary may require more seeds. (v) The constant  $c \approx 0.30$  is measured at a single parameter point ( $\alpha=0.06$ ,  $\sigma=0.06$ ,  $M=4$ ,  $\bar{K} \approx 8$ ,  $dt=0.02$ ). Its functional dependence on these parameters is not established here; back-of-envelope arguments for the scaling are given in Section 6 but require a systematic parameter sweep to validate. (vi) The  $c \approx 0.30$  result holds within the linear-rotation regime and is the significance-based threshold at two interior speeds. The slow condition exhibits a genuine sign reversal ( $L_{\text{max}} \approx L^*$ ); the mid condition’s  $L_{\text{max}}=3$  marks statistical detectability rather than coordination collapse ( $L^* \in (30, 45)$ ). The quantitative constant  $c \approx 0.30$  is primarily anchored on the slow condition, where the two operationalisations agree. (vii) The directional diagnostics (Figure 3) directly measure shared-arm angular bias and inter-agent dispersion at both interior speeds. They do not measure spatial trajectory overlap, spatial coverage, or provide a formal mediation chain from

angular dispersion to capture counts; that causal chain is consistent with the evidence but not directly tested. (viii) The discrepancy between the theoretical  $L^* \propto (\omega \cdot dt)^{-2/3}$  and the observed  $L_{\text{max}} \propto (\omega \cdot dt)^{-1}$  is attributed to policy nonlinearity, but the functional form of that nonlinearity is not derived or measured. This exponent discrepancy is an open problem. (ix) The four-regime taxonomy is data-derived, not pre-specified. A predictive rule for regime membership from  $(\omega, H)$  alone — the central open question of this work — has not been established.

## 7 Conclusion

We provided evidence for a stale-coordination failure mechanism in which the shared arm drives all collectors toward a systematically stale direction, reducing the inter-agent angular diversity that makes coordination beneficial. Instrumented runs directly show the mechanism’s two directional signatures (Figure 3); the causal chain from reduced angular diversity to reduced captures is supported by the performance crossover but not isolated by a formal mediation analysis. The coordination benefit landscape under rotation is *regime-dependent*: different rotation speeds produce qualitatively different phenomena, and identifying which regime applies is as important as the within-regime scaling constant.

Four regimes emerge from this study. *Linear-rotation regime* ( $\omega L_{\text{max}} dt \approx 0.30$  rad): a clean crossover at  $L_{\text{max}}/T_{\text{corr}} \approx 0.30$ ; empirical threshold earlier than the  $L^* \propto (\omega \cdot dt)^{-2/3}$  bias/noise prediction because the capacity-matching policy collapses coordination before directional parity is reached. *Sinc-attenuation regime* ( $\omega \in \{2.0, 2.5, 3.0\}$ ): amplitude damping scales down both the noise-reduction benefit and the stale-direction cost together, suppressing the sharp crossover. *Rapid-rotation regime* ( $\omega \in \{10.0, 17.0\}$ ):  $L_{\text{max}}$  falls below integer resolution; mechanism present but undetectable at  $n=32$ . *Long-episode regime* ( $\omega=1.0$ ,  $T_{\text{corr}} \approx H$ ): only  $L=1$  significant; mechanism unresolved.

Within the linear-rotation regime and the specific parameter setting studied ( $\alpha=0.06$ ,  $\sigma=0.06$ ,  $M=4$ ,  $dt=0.02$ , sliding-window controller), the heuristic  $L \lesssim 0.30/(\omega \cdot dt)$  avoids the stale-coordination failure. The  $c \approx 0.30$  constant rests on two interior anchors; the borderline slow anchor ( $p=0.049$ ) was confirmed at  $n=64$  in a pre-specified sensitivity analysis ( $t=1.87$ ,  $p<0.05$ ,  $df=63$ ). The central open question is what determines which regime a given  $(\omega, H)$  pair falls into — equivalently, what governs whether a sharp coordination crossover exists at all.

## A Full Per-Cell Results

Table 7 reports  $\bar{\Delta}$ , SE, and  $p$ -values for all window-method cells. Decay-method  $L_{\text{max}}$  results are summarised in Table 6 in the main text. Significance at  $p<0.05$  is marked with (\*).

## B Reproduction Gate Details

Gate run:  $\omega=0$ ,  $L=1$  (stateless sliding window), seeds 6001–6032.  $\bar{\Delta}=+1.188$  (SE = 0.213, 95% CI [+0.77, +1.61]); one-sided  $t$ -test:  $t=5.58$ ,  $df=31$ ,  $p<0.001$ ; sign count 20/32. This matches SPS-C03’s  $\bar{\Delta}=+1.19$  (SD 2.44, 20/32 positive) within rounding error, confirming the rotating-field implementation is correct at  $\omega=0$ . Note: the reproduction gate uses  $L=1$ , not  $L=H=67$ . The  $L=67$  stationary cell

**Table 7: Window method: all 40 primary non-gate cells plus  $\omega=10.0$  exploratory cells ( $n=32$  each).**

| $\omega$      | $L$ | $\bar{\Delta}$ | SE   | $p$   |
|---------------|-----|----------------|------|-------|
| very_slow     | L1  | +0.53          | 0.44 | 0.115 |
|               | L3  | +0.47          | 0.67 | 0.243 |
|               | L10 | +0.22          | 0.48 | 0.322 |
|               | L30 | +0.41          | 0.40 | 0.156 |
|               | L67 | +0.78*         | 0.46 | 0.044 |
| slow          | L1  | +1.47*         | 0.51 | 0.002 |
|               | L3  | +1.50*         | 0.41 | 0.000 |
|               | L5  | +1.66*         | 0.45 | 0.000 |
|               | L10 | +0.94*         | 0.57 | 0.049 |
|               | L20 | -0.38          | 0.47 | 0.788 |
|               | L30 | -0.28          | 0.52 | 0.705 |
|               | L45 | -0.47          | 0.59 | 0.786 |
| mid           | L67 | -0.44          | 0.63 | 0.758 |
|               | L1  | +0.88*         | 0.39 | 0.012 |
|               | L3  | +0.91*         | 0.34 | 0.004 |
|               | L5  | +0.22          | 0.41 | 0.296 |
|               | L10 | +0.25          | 0.54 | 0.320 |
|               | L20 | +0.56          | 0.46 | 0.113 |
|               | L30 | +0.16          | 0.50 | 0.378 |
| fast          | L45 | -0.25          | 0.41 | 0.727 |
|               | L67 | +0.03          | 0.43 | 0.471 |
|               | L1  | +0.47          | 0.44 | 0.145 |
|               | L3  | +0.47          | 0.44 | 0.145 |
|               | L5  | +0.78          | 0.39 | 0.023 |
|               | L10 | -0.09          | 0.52 | 0.572 |
|               | L20 | +0.09          | 0.46 | 0.420 |
| $\omega=10.0$ | L30 | -0.38          | 0.49 | 0.777 |
|               | L45 | +0.13          | 0.53 | 0.406 |
|               | L67 | +0.25          | 0.50 | 0.310 |
|               | L1  | +0.56          | 0.44 | 0.107 |
|               | L3  | +0.22          | 0.39 | 0.290 |
|               | L5  | +0.50          | 0.47 | 0.149 |
|               | L10 | +0.41          | 0.45 | 0.185 |
|               | L20 | +0.44          | 0.42 | 0.155 |
|               | L30 | +1.03*         | 0.51 | 0.025 |
|               | L45 | -0.19          | 0.63 | 0.616 |
|               | L67 | -0.16          | 0.55 | 0.612 |

in Table 2 is a separate full-history sliding-window cell that coincidentally yields  $\bar{\Delta}=+1.19$  because field stationarity makes window length irrelevant; it is *not* the gate cell.

### C Proof Sketch: SNR under Angular Bias

At step  $t$ , the window estimate  $\hat{\theta}(t, L)$  (Equation 1) is the mean of  $L$  observations at steps  $t, t-1, \dots, t-L+1$ . Its expected direction is  $\hat{\theta} = \theta(t) - \omega(L-1) dt/2$  (lag to the window midpoint) and its sinc-attenuated magnitude is  $\alpha \text{sinc}(\omega L dt/2)$  (the geometric-sum amplitude). Note that the direction lag uses  $(L-1)$  (the span from the oldest to the newest observation) while the sinc argument uses  $L$  (both  $\sin(L\omega dt/2)$  and the denominator  $L$  arise from the geometric sum, not the midpoint approximation). The direction formula and the sinc-attenuation approximation are valid only while the sinc factor is positive, i.e.  $\omega L dt/2 < \pi$  (equivalently

$L < \pi/(\omega dt)$ ). In the crossover region (cells at or below  $L_{\max}$ , where  $\omega L dt \leq 0.30$  rad) the argument is at most 0.15 rad, giving  $\text{sinc}(0.15) \approx 0.996$  ( $< 0.5\%$  amplitude loss), so sinc attenuation is negligible and the direction formula holds. Beyond  $L_{\max}$ , and in particular at very large  $L$ , the sinc factor can become negative: at  $(\omega=5, L=67)$  the argument is 3.35 rad and  $\text{sinc}(3.35) \approx -0.06$ . A negative sinc factor means the expected averaged vector *flips direction by  $\pi$* , so the directional formula no longer holds for such cells. No mechanistic claim is made here about the large- $L$  regime; it is outside the scope of this analysis, and those cells are empirically in the degraded region ( $\bar{\Delta} \leq 0$ ). The dominant factor in the crossover region is the angular lag  $\beta = \omega(L-1) dt/2$ .

The SNR of the direction estimate for  $M$  agents is approximately

$$\text{SNR}_{\text{shared}}(L, \omega) = \frac{\alpha \cos(\beta)}{\sigma/\sqrt{MKL} dt},$$

and the ratio to the independent arm is  $\sqrt{M}$ , independent of  $L$  and  $\omega$  — showing the simple model cannot predict negative  $\bar{\Delta}$ . Full derivation available on request.

### References

- [1] Anonymous Authors. 2026. *SPS-C03: Coordination Effect under Stationary Latent Field — Shared Summary Versus Independent Collectors*. Technical Report. Anonymous Institution. Preregistration and results available at [anonymized].
- [2] Liam Buffalo, John Pearson, and Andreas Klein. 2026. Realizing Common Random Numbers: Event-Keyed Hashing for Causally Valid Stochastic Models. *arXiv preprint arXiv:2603.11084* (2026). <https://arxiv.org/abs/2603.11084>
- [3] Abhishek Das, Théophile Gervet, Joshua Romoff, Dhruv Batra, Devi Parikh, Michael Rabbat, and Joelle Pineau. 2019. TarMAC: Targeted Multi-Agent Communication. In *International Conference on Machine Learning*. 1538–1546. <https://proceedings.mlr.press/v97/das19a.html>
- [4] Jakob N. Foerster, Ioannis A. Assael, Nando de Freitas, and Shimon Whiteson. 2016. Learning to Communicate with Deep Multi-Agent Reinforcement Learning. In *Advances in Neural Information Processing Systems*, Vol. 29. <https://proceedings.neurips.cc/paper/2016/hash/c7635bfd99248a2cdef8249ef7bfbef4-Abstract.html>
- [5] Aurélien Garivier and Eric Moulines. 2011. On Upper-Confidence Bound Policies for Non-Stationary Bandit Problems. *Algorithmic Learning Theory* (2011), 174–188. doi:10.1007/978-3-642-24412-4\_16
- [6] Paul Glasserman and David D. Yao. 1992. Some Guidelines and Guarantees for Common Random Numbers. *Management Science* 38, 6 (1992), 884–908. doi:10.1287/mnsc.38.6.884
- [7] Steven Kapturowski, Georg Ostrovski, John Quan, Remi Munos, and Will Dabney. 2019. Recurrent Experience Replay in Distributed Reinforcement Learning. In *International Conference on Learning Representations*. <https://openreview.net/forum?id=r1lyTjAqYX>
- [8] Levente Kocsis and Csaba Szepesvári. 2006. Discounted UCB. In *2nd PASCAL Challenges Workshop*. 784–791.
- [9] David Patiño, Siddharth Mayya, Juan A. Calderon, Kostas Daniilidis, and David Saldaña. 2023. Learning to Navigate in Turbulent Flows with Aerial Robot Swarms: A Cooperative Deep Reinforcement Learning Approach. In *arXiv preprint arXiv:2306.04781*. <https://arxiv.org/abs/2306.04781>
- [10] Amanpreet Singh, Tushar Jain, and Sainbayar Sukhbaatar. 2019. Learning when to Communicate at Scale in Multiagent Cooperative and Competitive Tasks. In *International Conference on Learning Representations*. <https://openreview.net/forum?id=rye7knCqK7>
- [11] Fangyuan Wang, Julio M. Ottino, Paul B. Umbanhowar, and Richard M. Lueptow. 2025. Mobile-collector capture of particles in a chaotic flow. *PLOS ONE* (2025). doi:10.1371/journal.pone.0329766
- [12] Tianhao Wang, Jiachen Dong, Victor Lesser, and Chongjie Zhang. 2020. Learning Efficient Multi-agent Communication: An Information Bottleneck Approach. In *International Conference on Machine Learning*. 9908–9918. <https://proceedings.mlr.press/v119/wang20i.html>