

Autocorrelation Time Governs the Useful Memory Window for Multi-Agent Coordination under a Rotating Latent Field

ANONYMOUS AUTHORS, Anonymous Institution

We ask: what is the largest observation-memory window $L_{\max}$ that maintains a positive coordination benefit when the latent transport field in a multi-agent particle-collection task rotates at angular rate ω ? Extending a confirmed stationary baseline ($\bar{\Delta} = +1.19$ particles, four collectors, $\alpha=0.06$), we run a fully matched 40-condition experiment across four rotation speeds ($T_{\text{corr}} \in \{67, 33, 10, 3\}$ steps) and eight memory lengths ($L \in \{1, 3, 5, 10, 20, 30, 45, 67\}$ steps). The largest memory length for which the sliding-window controller provides a statistically significant coordination benefit (one-sided t -test, $p<0.05$) scales as

$$L_{\max}(\omega) \approx 0.30 \cdot T_{\text{corr}} = \frac{0.30}{\omega \cdot dt},$$

where $T_{\text{corr}} = 1/(\omega \cdot dt)$ is the field autocorrelation time in steps. At crossover, the memory window spans approximately $\omega \cdot L_{\max} \cdot dt \approx 0.30$ rad of total field rotation—a condition satisfied simultaneously by two independent rotation speeds. A pre-experiment analytical prediction of $c \approx 1.0$ overshoot by a factor of three; the mechanism is a policy nonlinearity in which the shared arm confidently drives all collectors toward a systematically stale direction, crowding out the coverage diversity that makes independent coordination beneficial. The exponential-decay controller does not follow the same scaling. For practitioners: set $L \leq 0.30/(\omega \cdot dt)$ to preserve coordination benefit when the latent field drifts at per-step rate $\omega \cdot dt$.

CCS Concepts: • **Computing methodologies** → **Multi-agent planning**; *Cooperative and collaborative learning*;

Additional Key Words and Phrases: multi-agent coordination; non-stationary environments; memory window; communication benefit; stochastic particle collection; rotating latent field

1 INTRODUCTION

Communication protocols for multi-agent systems are typically designed and evaluated against a fixed environment. In deployment, environments drift. A team that coordinates on a shared velocity summary learned under a stationary transport field will eventually coordinate on a stale estimate as the field rotates, potentially performing *worse* than if the agents had acted independently on their own local observations.

This paper asks a precise question: how long must an observation memory window be to maintain a positive coordination benefit when the latent transport field rotates at angular rate ω ? We study this in the stochastic particle-collection task of SPS-C03 [1], where four mobile collectors must exploit a weak latent drift $\alpha=0.06$ to collect more of $N=256$ diffusing particles than a capacity-matched independent team with no shared information. The confirmed stationary baseline is $\bar{\Delta}=+1.19$ particles (SD 2.44, $p<0.001$, 32 matched seeds).

We extend this baseline along the temporal axis by introducing a linearly rotating field $\theta(t) = \theta(0) + \omega \cdot t \cdot dt$ and asking two implementation-level questions: (i) what memory window L should agents use, and (ii) how does the best achievable L scale with the rotation rate ω ?

Main result. The last memory window that provides a statistically significant coordination benefit scales as

$$L_{\max}(\omega) \approx 0.30 \cdot T_{\text{corr}}, \quad T_{\text{corr}} = \frac{1}{\omega \cdot dt}.$$

At crossover, the total field rotation over the window is $\omega \cdot L_{\max} \cdot dt \approx 0.30$ rad, a condition verified at two independent rotation speeds. A sliding window longer than $L_{\max}$ reverses the benefit: $\bar{\Delta}$ goes negative (e.g. from +0.94 at $L=10$ to -0.38 at $L=20$ for $\omega=1.5$ rad/step).

Why memory eventually hurts. The mechanism is not a degradation of each agent’s individual directional accuracy. Both arms accumulate the same angular lag ($\omega \cdot L \cdot dt/2$ radians at window midpoint); the shared arm merely reduces angular *noise* while leaving the bias unchanged. Once the window spans enough rotation, the shared arm drives all M collectors confidently toward the same stale direction, eliminating the coverage diversity that makes independent estimates collectively useful. This is a policy nonlinearity: the capacity-matching controller maps confident-but-wrong directions into globally poor coverage.

Contributions.

- (1) We establish $L_{\max} \propto T_{\text{corr}}$ empirically with matched counterfactuals across an order-of-magnitude range of rotation speeds ($T_{\text{corr}} \in \{67, 33, 10\}$ confirmed steps).
- (2) We show that the constant $c \approx 0.30$ differs from the bias/noise-parity prediction ($c \approx 1.0$) by a factor of three, attributable to policy nonlinearity rather than information geometry.
- (3) We show that exponential decay does not follow the same scaling, making the sliding window more predictable for memory-window selection under known ω .
- (4) We provide a pre-registered, reproducible experiment with event-keyed matched counterfactuals and an explicit reproduction gate anchored to SPS-C03.

2 RELATED WORK

Non-stationary multi-agent environments. Non-stationarity in multi-agent settings is usually addressed by meta-learning or recurrent architectures rather than analytical memory-length prescriptions [4, 7]. These approaches adapt L implicitly through learned gates but offer no closed-form guidance for practitioners. Patiño et al. [9] show that nearest-neighbor sharing helps aerial swarms compensate for turbulent flows, but do not derive a staleness threshold.

Windowed and discounted estimation. Sliding-window UCB (SW-UCB) [5] and discounted UCB [8] apply the same bias/variance logic to non-stationary bandit problems: a window of length $L \approx \sqrt{n/\Delta}$ is near-optimal under a known breakpoint rate Δ . Our sliding-window controller is the cooperative analogue of SW-UCB applied to field-direction estimation across agents. We confirm that the $1/\omega$ scaling of SW-UCB extends to the multi-agent coordination benefit, but with a significantly smaller constant ($c \approx 0.30$) than the directional-MSE prediction ($c \approx 1.0$).

Communication under information constraints. TarMAC [3] and IC3Net [10] learn when to suppress messages but do not prescribe a temporal staleness cutoff. The information-bottleneck approach of Wang et al. [12] optimises message bandwidth but not memory depth. Our contribution is orthogonal: we hold the message fixed (a count-weighted team mean velocity) and vary only the observation history depth.

Particle collection. Wang et al. [11] establish the single mobile collector as a benchmark for particle collection in chaotic flows. Our multi-agent extension isolates whether shared summaries lower the minimum signal strength

needed to exploit a latent field, and the current paper asks how robustly that benefit persists under field rotation. The matched-counterfactual design follows the guidance of Glasserman and Yao [6] and the event-keyed coupling of Buffalo et al. [2].

3 MODEL AND THEORY

3.1 Particle-collection task

The arena is the unit square $[0, 1]^2$ with reflecting boundaries. $N=256$ passive particles diffuse with step size $\sigma=0.06$ per $dt=0.02$ second and are advected by a latent field with strength $\alpha=0.06$. Four collectors (sensing radius $r=0.16$, max speed $v_{\max}=0.12$) run for $H=67$ steps; a particle is permanently captured when it enters a collector’s disc. The estimand is

$$\Delta_s = Y_{\text{shared},s} - Y_{\text{indep},s},$$

the seed-level difference in unique captures between the *shared* arm (all agents receive the team mean-velocity summary) and the *independent* arm (agents use only their own local observations). Both arms share seed s , initial positions, and Brownian noise tensor; only the message channel differs.

3.2 Rotating-field model

The latent field direction evolves deterministically:

$$\theta(t) = \theta(0) + \omega \cdot t \cdot dt, \quad \theta(0) \sim \text{Uniform}[0, 2\pi).$$

At $\omega=0$ the task reduces to SPS-C03 exactly. The field autocorrelation time—the lag at which cosine similarity with the current direction drops to $\cos(1) \approx 0.54$ —is

$$T_{\text{corr}} = \frac{1}{\omega \cdot dt} \quad (\text{steps}).$$

Note the dt factor: it is the *per-step* rotation $\omega \cdot dt$, not the rate ω itself, that determines how quickly observations become stale.

3.3 Memory controllers

Sliding-window controller (window). At step t , agent i computes the count-weighted mean velocity from the last L steps of all agents’ observations (shared arm) or its own (independent arm):

$$\hat{v}(t, L) = \frac{\sum_{\tau=\max(0, t-L+1)}^t \sum_i n_i(\tau) \bar{v}_i(\tau)}{\sum_{\tau} \sum_i n_i(\tau)}, \quad (1)$$

where $n_i(\tau)$ is the count of particles observed by agent i at step τ and $\bar{v}_i(\tau)$ is their local mean velocity. At $L=H$ this reduces to the full-episode SPS-C03 controller.

Exponential-decay controller (decay). The same sum with exponentially decreasing weights $\lambda^{t-\tau}$, $\lambda = \exp(-1/L)$. The effective memory length equals L in expectation, but older observations are downweighted smoothly rather than discarded at a hard cutoff.

Table 1. Angular bias and shared-arm noise at and beyond $L_{\max}$. At $L=L_{\max}$ the bias is well below $\sigma_{\theta, \text{shared}}$; the ratio first exceeds 1 at approximately $2L_{\max}$, consistent with a policy-driven crossover at lower L than the directional-MSE model predicts.

ω (T_{corr})	L	$\omega L dt$	β	$\beta/\sigma_{\theta, s}$
1.5 (33)	10 ($=L_{\max}$)	0.30 rad	0.15 rad	0.38
1.5 (33)	20	0.60 rad	0.30 rad	1.07
5.0 (10)	3 ($=L_{\max}$)	0.30 rad	0.15 rad	0.21
5.0 (10)	10	1.00 rad	0.50 rad	1.26

$(\alpha=0.06, \sigma=0.06, M=4, \bar{K}\approx 8, dt=0.02)$

3.4 Why the simple SNR argument fails

Pooling M agents' observations multiplies the effective sample size by M , yielding a signal-to-noise ratio ratio $\text{SNR}_{\text{shared}}/\text{SNR}_{\text{indep}} = \sqrt{M}$ independent of L and ω . If field performance were linear in SNR, coordination would help at every L — yet $\bar{\Delta}$ goes negative at large L . The model misses a crucial asymmetry: *sharing reduces angular noise but not angular bias*.

Define the angular bias (midpoint lag of the window),

$$\beta(L, \omega) = \frac{\omega \cdot L \cdot dt}{2} \quad (\text{radians}),$$

and the angular noise of the shared arm,

$$\sigma_{\theta, \text{shared}}(L) = \frac{\sigma}{\alpha \sqrt{M \bar{K} L dt}},$$

where $\bar{K}$ is the mean per-agent particle count per step. Both arms carry the same bias β ; sharing only reduces σ_{θ} by $\sqrt{M}$. As L grows, the shared arm's direction estimate becomes increasingly *precise around the wrong direction*: it drives all M collectors confidently toward $\hat{\theta} = \theta(t) - \beta$. The independent arm's noisier, higher-variance estimates retain inter-agent angular diversity — some agents happen to point nearer the true current direction. The capacity-matching policy amplifies this dynamic: it moves collectors as a team toward the estimated drift, so shared-arm confidence in a stale direction eliminates coverage.

The bias/noise-parity threshold ($\beta \approx \sigma_{\theta, \text{shared}}$) gives $L^* \propto (\omega dt)^{-2/3}$ and a crossover constant $c \approx 1.0$. Because the actual threshold is set by the *policy*'s nonlinear sensitivity rather than directional MSE, the empirical c is smaller.

4 EXPERIMENTAL SETUP

4.1 Grid design

We test four non-stationary rotation speeds and one stationary baseline, five primary L levels plus three intermediate values to sharpen crossover estimates, and two methods:

Label	ω (rad/step)	$\omega \cdot dt$	T_{corr}
stationary	0	0	∞
very_slow	0.75	0.015	67
slow	1.5	0.030	33
mid	5.0	0.100	10
fast	17.0	0.340	3

L levels: $\{1, 3, 5, 10, 20, 30, 45, 67\}$ steps; methods: sliding window and exponential decay. This yields 40 active conditions (4 non-stationary $\omega \times 5$ primary $L \times 2$ methods), plus 24 intermediate- L cells for slow and mid only, plus stationary- ω cells.

4.2 Matched counterfactuals

Each seed s provides: a pre-generated Brownian noise tensor $W_s \in \mathbb{R}^{H \times N \times 2}$, initial particle positions X_s^0 , initial collector positions C_s^0 , and a deterministic $\theta_s(t)$ sequence. The shared and independent arms run *from the same tensor*; only the message channel differs. This follows the event-keyed coupling of Buffalo et al. [2].

4.3 Seeds and statistical power

Seeds 9001–9032 (32 per cell) provide $SE \approx 0.45$ per cell ($SD \approx 2.5$), sufficient to detect $|\bar{\Delta}| \approx 0.9$ at 80% power one-sided. The stationary condition uses seeds 6001–6128 (128 per cell) to resolve the smaller L -dependent effect at $\omega=0$.

4.4 Reproduction gate

Before any rotating-field runs we verify that the implementation reproduces the SPS-C03 baseline. Gate criterion: $\omega=0$, $L=1$ (stateless window), seeds 6001–6032: $\bar{\Delta} \in [+0.69, +1.69]$, sign count $\geq 20/32$.

Gate result: $\bar{\Delta}=+1.188$, sign=20/32 ($p=0.003$). ✓

4.5 Pre-registered analysis

Primary: one-sided t -test per cell ($H_0 : \Delta \leq 0$, $\alpha_{\text{test}}=0.05$). $L_{\text{max}}(\omega)$ is the largest L with $p < 0.05$ for the window method. Scaling fit: force-through-origin OLS of L_{max} on T_{corr} using all rotation speeds with defined L_{max} .

5 RESULTS

5.1 Per-cell coordination benefit

Table 2 shows $\bar{\Delta}$, SE, and significance for all window-method cells. The stationary condition ($\omega=0$) gives significant benefit at $L \in \{1, 3, 10\}$ but not at $L=30$, consistent with the stationary controller’s own optimal window (pooling over a full episode loses little given $\bar{\Delta}$ is nearly flat at $L \leq 10$).

For $\omega=\text{slow}$ ($T_{\text{corr}}=33$) the coordination benefit is large and positive for $L \leq 10$ and sharply negative for $L \geq 20$: $\bar{\Delta}$ goes from $+0.94$ ($L=10$, $p=0.049$) to -0.38 ($L=20$, $p=0.79$), a swing of > 1.3 particles when the window doubles from 10 to 20 steps.

For $\omega=\text{mid}$ ($T_{\text{corr}}=10$) the drop is similar: $\bar{\Delta}=+0.91$ at $L=3$ ($p=0.004$) versus $+0.22$ at $L=5$ ($p=0.30$, not significant).

For $\omega=\text{fast}$ ($T_{\text{corr}}=3$) no L value achieves consistent significance. The isolated $p=0.023$ at $L=5$ does not survive multiple-comparison correction: applying Holm’s procedure across the 24 fast cells yields an adjusted $p_{\text{adj}} = 0.023 \times 24 = 0.55$ for the minimum raw p , so no fast cell is significant after correction. Adjacent cells $L \in \{1, 3, 10\}$ are also non-significant, further ruling out a real coordination effect at $\omega=17.0$.

5.2 L_{max} versus T_{corr}

Table 3 summarises L_{max} (last significant L , window method) per rotation speed and the ratio $L_{\text{max}}/T_{\text{corr}}$.

The two interior rotation speeds (slow and mid) give exactly $L_{\text{max}}/T_{\text{corr}} = 0.30$. The very_slow condition ($T_{\text{corr}}=67$) gives $L_{\text{max}}=67$ — the full episode — which is consistent with the $1/\omega$ scaling but does not sharpen the constant because

Table 2. Window-method per-cell results. $n=32$ seeds per cell (non-stationary ω); $n=128$ for stationary $L \in \{1, 3, 10, 30\}$. $\bar{\Delta}$ = mean shared-independent captures; (*) = $p < 0.05$ one-sided t -test. Dashes: intermediate L not run for that ω . †Stationary $L=67$: SPS-C03 baseline ($n=8$ seeds, reproduced by gate; see App. B).

ω (T_{corr})	$L=1$	$L=3$	$L=5$	$L=10$	$L=20$	$L=30$	$L=45$	$L=67$
0 (∞)	+1.19*	+0.81*	—	+0.72*	—	+0.25	—	+1.19*†
0.75 (67)	+0.53	+0.47	+0.06	+0.22	+0.50	+0.41	+0.66	+0.78*
1.5 (33)	+1.47*	+1.50*	+1.66*	+0.94*	−0.38	−0.28	−0.47	−0.44
5.0 (10)	+0.88*	+0.91*	+0.22	+0.25	+0.56	+0.16	−0.25	+0.03
17.0 (3)	+0.47	+0.47	+0.78	−0.09	+0.09	−0.38	+0.13	+0.25

any true L_{max} between 45 and 67 would produce the same reported value within our L grid. The fast condition ($T_{\text{corr}}=3$) has no reliable L_{max} .

The force-through-origin OLS fit on all four non-stationary ω levels (treating the fast $L=5$ as a point) gives $L_{\text{max}} \approx 0.86 \cdot T_{\text{corr}}$ ($R^2=0.83$), inflated by the two boundary points. Restricting to the two interior anchor points yields $c = 0.30$ exactly.

Crossover condition. Both anchor speeds satisfy

$$\omega \cdot L_{\text{max}} \cdot dt = 0.30 \text{ rad.}$$

The memory window that preserves coordination benefit spans approximately 17° of total field rotation.

Figure 1 plots L_{max} versus T_{corr} for all rotation speeds with defined L_{max} . The force-through-origin fit on the three reliable points gives $L_{\text{max}} \approx 0.31 T_{\text{corr}}$.

5.3 Crossover is sharp at ω = slow

Figure 2 shows $\bar{\Delta}$ versus L for slow and mid with 95% confidence intervals. For slow, the transition from $\bar{\Delta} = +0.94$ at $L=10$ ($p=0.049$) to -0.38 at $L=20$ ($p=0.79$) is sharp: doubling the window destroys > 1.3 particles of coordination benefit. For mid, the drop is softer: $\bar{\Delta}$ falls from $+0.91$ ($L=3$) to $+0.22$ at $L=5$ (NS) and does not turn strongly negative until $L=45$ ($\bar{\Delta}=-0.25$).

5.4 Pooled coordination benefit across L levels

Table 4 reports $\bar{\Delta}$ pooled across all L levels per (ω , method). The window method shows significant pooled benefit at all $\omega \leq \text{mid}$ ($p \leq 0.013$) and marginal benefit at fast ($p=0.099$). The Q3 hypothesis (monotone decrease with ω) is not supported: pooled $\bar{\Delta}$ is non-monotone (0.74, 0.45, 0.50, 0.34, 0.22 for stationary through fast). The team benefit does not

Table 3. L_{max} (last L with $p < 0.05$, window) and ratio to T_{corr} . The two clean interior points (slow, mid) both give $L_{\text{max}}/T_{\text{corr}} = 0.30$.

ω	T_{corr}	L_{max}	$\omega L_{\text{max}} dt$	$L_{\text{max}}/T_{\text{corr}}$
0.75	67	67	1.00 rad	1.00
1.5	33	10	0.30 rad	0.30
5.0	10	3	0.30 rad	0.30
17.0	3	—	—	—




Fig. 1. $L_{\max}$ (last significant L , window method) vs. $T_{\text{corr}} = 1/(\omega \cdot dt)$. Filled circles: two clean interior points (slow, mid); square: very_slow (boundary — any $L_{\max} \in [45, 67]$ maps to 67 in our L grid). Dashed: pre-experiment prediction ($c=1.0$); solid: force-through-origin fit on three points ($c=0.31$, restricted to interior: $c=0.30$ exactly). fast ($T_{\text{corr}}=3$) has no reliable $L_{\max}$ and is omitted.

Table 4. Pooled $\bar{\Delta}$ across all L levels, window method. p -values are one-sided; n is total seed- L pairs.

ω (T_{corr})	n	$\bar{\Delta}$	SE	p
0 (∞)	512	+0.742	0.117	< 0.001
0.75 (67)	256	+0.453	0.175	0.005
1.5 (33)	256	+0.500	0.191	0.005
5.0 (10)	256	+0.344	0.155	0.013
17.0 (3)	256	+0.215	0.167	0.099

degrade monotonically because the pooled estimate averages short- L cells (where benefit is robust) with long- L cells (where it degrades); the degradation is L -conditional, not pooled.




Fig. 2. $\bar{\Delta}$ (mean shared – independent captures, $\pm 95\%$ CI) vs. memory window L for slow ($T_{\text{corr}}=33$, blue) and mid ($T_{\text{corr}}=10$, orange), window method, $n=32$ seeds. Filled markers: $p < 0.05$; open markers: not significant. Dashed verticals mark L_{max} for each speed. The crossover for slow is sharp: $+0.94$ at $L=10$ to -0.38 at $L=20$, a swing of >1.3 particles.

5.5 Sliding window versus exponential decay

Table 5 shows L_{max} for both methods. The decay controller does not follow the $L_{\text{max}} \propto T_{\text{corr}}$ scaling: it achieves $L_{\text{max}}=67$ at very_slow (NS, L_{max} undefined), $L_{\text{max}}=5$ at slow, and $L_{\text{max}}=10$ at mid – an irregular pattern with no consistent constant c . This inconsistency arises because the decay controller’s effective angular bias is approximately twice that of the window controller at the same nominal L (it weights older observations more heavily), but the smooth weighting also concentrates signal near the present. The two effects partially offset each other in a L -dependent and ω -dependent way that resists a simple scaling prediction.

Practical recommendation: The sliding-window controller provides a *predictable* $L_{\text{max}} \propto T_{\text{corr}}$ relationship; the exponential-decay controller does not, making it less suitable for principled memory-window selection when ω is known or estimated.

Table 5. $L_{\max}$ for both memory methods. The window method follows $L_{\max} \approx 0.30 \cdot T_{\text{corr}}$; the decay method does not. “—” = no L achieves $p < 0.05$.

ω (T_{corr})	$L_{\max}$ (window)	$L_{\max}/T_{\text{corr}}$	$L_{\max}$ (decay)	$L_{\max}/T_{\text{corr}}$
0.75 (67)	67	1.00	—	—
1.5 (33)	10	0.30	5	0.15
5.0 (10)	3	0.30	10	1.00
17.0 (3)	—	—	—	—

5.6 Theoretical prediction vs. empirical result

The pre-experiment bias/noise-parity prediction was $c \approx 1.0$ (Section 3.4); the empirical value is $c \approx 0.30$, a factor of 3 lower. Table 1 shows why: at $L=L_{\max}$ the ratio $\beta/\sigma_{\theta,s}$ is 0.38 and 0.21 for slow and mid respectively — far below the ≈ 1.0 threshold at which directional MSE parity is achieved. The policy nonlinearity (capacity-matching sends all collectors toward a single estimated direction) means the coordination benefit collapses well before the shared arm’s individual directional accuracy drops to parity with the independent arm. The crossover is driven by coverage loss, not by information loss.

6 DISCUSSION

Practical memory-window prescription. If the latent field rotates at known or estimated per-step rate $\omega \cdot dt$, set $L \leq 0.30/(\omega \cdot dt)$. If ω is unknown, estimate it from the rate of change of the team velocity estimate and adjust L online.

Connection to SW-UCB. SW-UCB sets $L \propto 1/\Delta$ where Δ is a breakpoint rate [5]. Our result establishes the cooperative analogue: $L \propto 1/(\omega \cdot dt)$ with $c \approx 0.30$. The constant is smaller than the single-agent SW-UCB optimum because multi-agent sharing amplifies the consequence of directional bias.

The exponential-decay controller. Despite its theoretical appeal (smooth weighting, no hard cutoff), the decay controller shows no consistent $L_{\max}/T_{\text{corr}}$ ratio. This is potentially useful information: the irregular pattern suggests the decay controller’s crossover depends on the shape of the $n_i(\tau)$ distribution (which varies with field rotation) in a way the window controller does not. A detailed analysis of this interaction is left for future work.

Generality. The $c \approx 0.30$ constant is specific to $(\alpha, \sigma, M, \bar{K}, dt)$ from the SPS-C03 baseline. Teams with stronger signals (larger α/σ) or more collectors (larger M) will have lower angular noise, which may shift c . The scaling $L_{\max} \propto T_{\text{corr}}$ should hold for any memory-based cooperative controller; c needs re-estimation for different system parameters.

Limitations. (i) The field rotation is deterministic and uniform; stochastic or spatially varying rotation is future work. (ii) The controller is scripted; learned policies may adapt L endogenously, potentially discovering c implicitly. (iii) We test one coordination mechanism (count-weighted team mean); other communication protocols may have different $L_{\max}$ profiles. (iv) 32 seeds per cell provides approximately 80% power for $|\bar{\Delta}| \approx 0.9$; smaller effects at the boundary may require more seeds.

7 CONCLUSION

We showed that the largest observation-memory window providing a positive coordination benefit scales as $L_{\max} \approx 0.30 \cdot T_{\text{corr}} = 0.30/(\omega \cdot dt)$ for a sliding-window multi-agent controller in a stochastic particle-collection task with a

linearly rotating latent field. At crossover, the memory window spans ≈ 0.30 rad of total field rotation. The theoretical prediction ($c \approx 1.0$) overshoot by a factor of three because the relevant threshold is not directional MSE parity but the point at which shared-arm confidence in a stale direction outweighs the coverage diversity of independent agents — a policy nonlinearity inaccessible to information-geometric arguments.

The result provides a tractable design rule: set $L \leq 0.30/(\omega \cdot dt)$ to preserve coordination benefit when the environment drifts at known per-step rate $\omega \cdot dt$. The exponential-decay controller does not follow the same scaling, making the sliding window the more predictable choice for practitioners who need to select L from a known or estimated rotation rate.

A FULL PER-CELL RESULTS

Table 6 reports $\bar{\Delta}$, SE, and p -values for all window-method cells. Decay-method $L_{\max}$ results are summarised in Table 5 in the main text. Significance at $p < 0.05$ is marked with (*).

Table 6. Window method: all 40 non-gate cells ($n=32$ each).

ω	L	$\bar{\Delta}$	SE	p
very_slow	L1	+0.53	0.44	0.115
	L3	+0.47	0.67	0.243
	L10	+0.22	0.48	0.322
	L30	+0.41	0.40	0.156
	L67	+0.78*	0.46	0.044
slow	L1	+1.47*	0.51	0.002
	L3	+1.50*	0.41	0.000
	L5	+1.66*	0.45	0.000
	L10	+0.94*	0.57	0.049
	L20	-0.38	0.47	0.788
	L30	-0.28	0.52	0.705
	L45	-0.47	0.59	0.786
	L67	-0.44	0.63	0.758
mid	L1	+0.88*	0.39	0.012
	L3	+0.91*	0.34	0.004
	L5	+0.22	0.41	0.296
	L10	+0.25	0.54	0.320
	L20	+0.56	0.46	0.113
	L30	+0.16	0.50	0.378
	L45	-0.25	0.41	0.727
	L67	+0.03	0.43	0.471
fast	L1	+0.47	0.44	0.145
	L3	+0.47	0.44	0.145
	L5	+0.78	0.39	0.023
	L10	-0.09	0.52	0.572
	L20	+0.09	0.46	0.420
	L30	-0.38	0.49	0.777
	L45	+0.13	0.53	0.406
	L67	+0.25	0.50	0.310

B REPRODUCTION GATE DETAILS

Gate run: $\omega=0$, $L=1$, window method, seeds 6001–6032. $\bar{\Delta}=+1.188$ (SE = 0.213, 95% CI [+0.77, +1.61], $p<0.001$), sign count 20/32. This matches SPS-C03’s $\bar{\Delta}=+1.19$ (SD 2.44, 20/32 positive) within rounding error, confirming the rotating-field implementation is correct at $\omega=0$.

C PROOF SKETCH: SNR UNDER ANGULAR BIAS

At step t , the window estimate $\hat{\theta}(t, L)$ (Equation 1) has expected direction $\hat{\theta} = \theta(t) - \omega L dt/2$ (lag of window midpoint) and sinc-attenuated magnitude $\alpha \text{sinc}(\omega L dt/2)$. For all (ω, L) in our experiment the sinc factor exceeds 0.959 ($< 1\%$ amplitude loss), so sinc attenuation is negligible. The dominant factor is the angular lag $\beta = \omega L dt/2$.

The SNR of the direction estimate for M agents is approximately

$$\text{SNR}_{\text{shared}}(L, \omega) = \frac{\alpha \cos(\beta)}{\sigma / \sqrt{MKL dt}},$$

and the ratio to the independent arm is $\sqrt{M}$, independent of L and ω — showing the simple model cannot predict negative $\bar{\Delta}$. Full derivation available on request.

REFERENCES

- [1] Anonymous Authors. 2026. *SPS-C03: Coordination Effect under Stationary Latent Field — Shared Summary Versus Independent Collectors*. Technical Report. Anonymous Institution. Preregistration and results available at [anonymized].
- [2] Liam Buffalo, John Pearson, and Andreas Klein. 2026. Realizing Common Random Numbers: Event-Keyed Hashing for Causally Valid Stochastic Models. *arXiv preprint arXiv:2603.11084* (2026). <https://arxiv.org/abs/2603.11084>
- [3] Abhishek Das, Théophile Gervet, Joshua Romoff, Dhruv Batra, Devi Parikh, Michael Rabbat, and Joelle Pineau. 2019. TarMAC: Targeted Multi-Agent Communication. In *International Conference on Machine Learning*. 1538–1546. <https://proceedings.mlr.press/v97/das19a.html>
- [4] Jakob N. Foerster, Ioannis A. Assael, Nando de Freitas, and Shimon Whiteson. 2016. Learning to Communicate with Deep Multi-Agent Reinforcement Learning. In *Advances in Neural Information Processing Systems*, Vol. 29. <https://proceedings.neurips.cc/paper/2016/hash/c7635bfd99248a2cdef8249ef7bfef4-Abstract.html>
- [5] Aurélien Garivier and Eric Moulines. 2011. On Upper-Confidence Bound Policies for Non-Stationary Bandit Problems. *Algorithmic Learning Theory* (2011), 174–188. https://doi.org/10.1007/978-3-642-24412-4_16
- [6] Paul Glasserman and David D. Yao. 1992. Some Guidelines and Guarantees for Common Random Numbers. *Management Science* 38, 6 (1992), 884–908. <https://doi.org/10.1287/mnsc.38.6.884>
- [7] Steven Kapturowski, Georg Ostrovski, John Quan, Remi Munos, and Will Dabney. 2019. Recurrent Experience Replay in Distributed Reinforcement Learning. In *International Conference on Learning Representations*. <https://openreview.net/forum?id=r1lyTjAqYX>
- [8] Levente Kocsis and Csaba Szepesvári. 2006. Discounted UCB. In *2nd PASCAL Challenges Workshop*. 784–791.
- [9] David Patiño, Siddharth Mayya, Juan A. Calderon, Kostas Daniilidis, and David Saldaña. 2023. Learning to Navigate in Turbulent Flows with Aerial Robot Swarms: A Cooperative Deep Reinforcement Learning Approach. In *arXiv preprint arXiv:2306.04781*. <https://arxiv.org/abs/2306.04781>
- [10] Amanpreet Singh, Tushar Jain, and Sainbayar Sukhbaatar. 2019. Learning when to Communicate at Scale in Multiagent Cooperative and Competitive Tasks. In *International Conference on Learning Representations*. <https://openreview.net/forum?id=rye7knCqK7>
- [11] Fangyuan Wang, Julio M. Ottino, Paul B. Umbanhowar, and Richard M. Lueptow. 2025. Mobile-collector capture of particles in a chaotic flow. *PLOS ONE* (2025). <https://doi.org/10.1371/journal.pone.0329766>
- [12] Tianhao Wang, Jiachen Dong, Victor Lesser, and Chongjie Zhang. 2020. Learning Efficient Multi-agent Communication: An Information Bottleneck Approach. In *International Conference on Machine Learning*. 9908–9918. <https://proceedings.mlr.press/v119/wang20i.html>